



>>> IMAGE PROCESSING AND  
COMPUTATIONAL PHOTOGRAPHY  
SESSION 10: HOMOGRAPHIES, PANORAMAS  
AND PERSPECTIVE

Oriol Pujol & Simone Balocco

Inspired by D. Hoiem slides

# Why Mosaic?

Are you getting the whole picture?

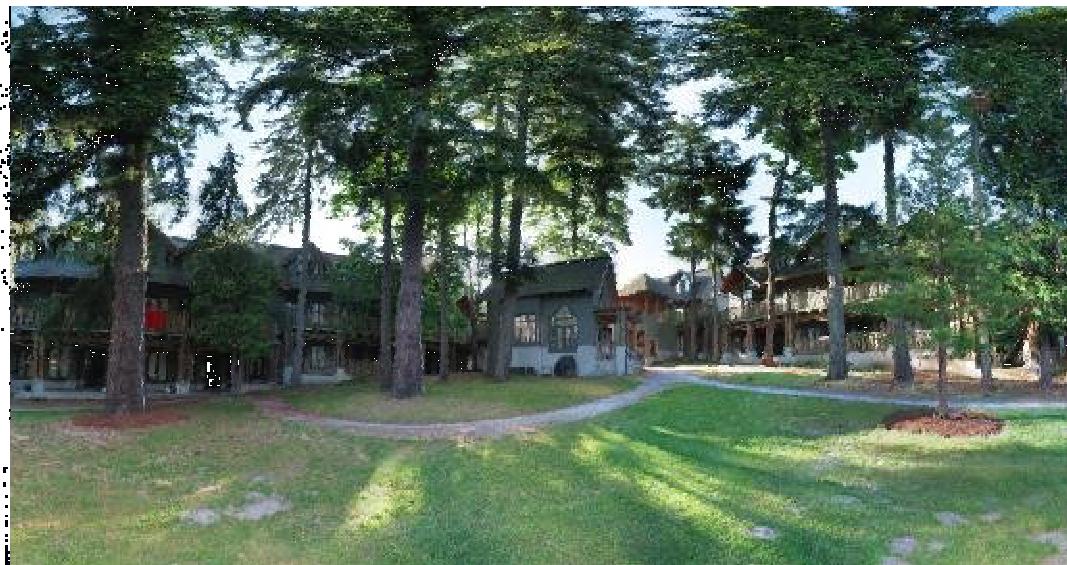
- Compact Camera FOV =  $50 \times 35^\circ$



# Why Mosaic?

Are you getting the whole picture?

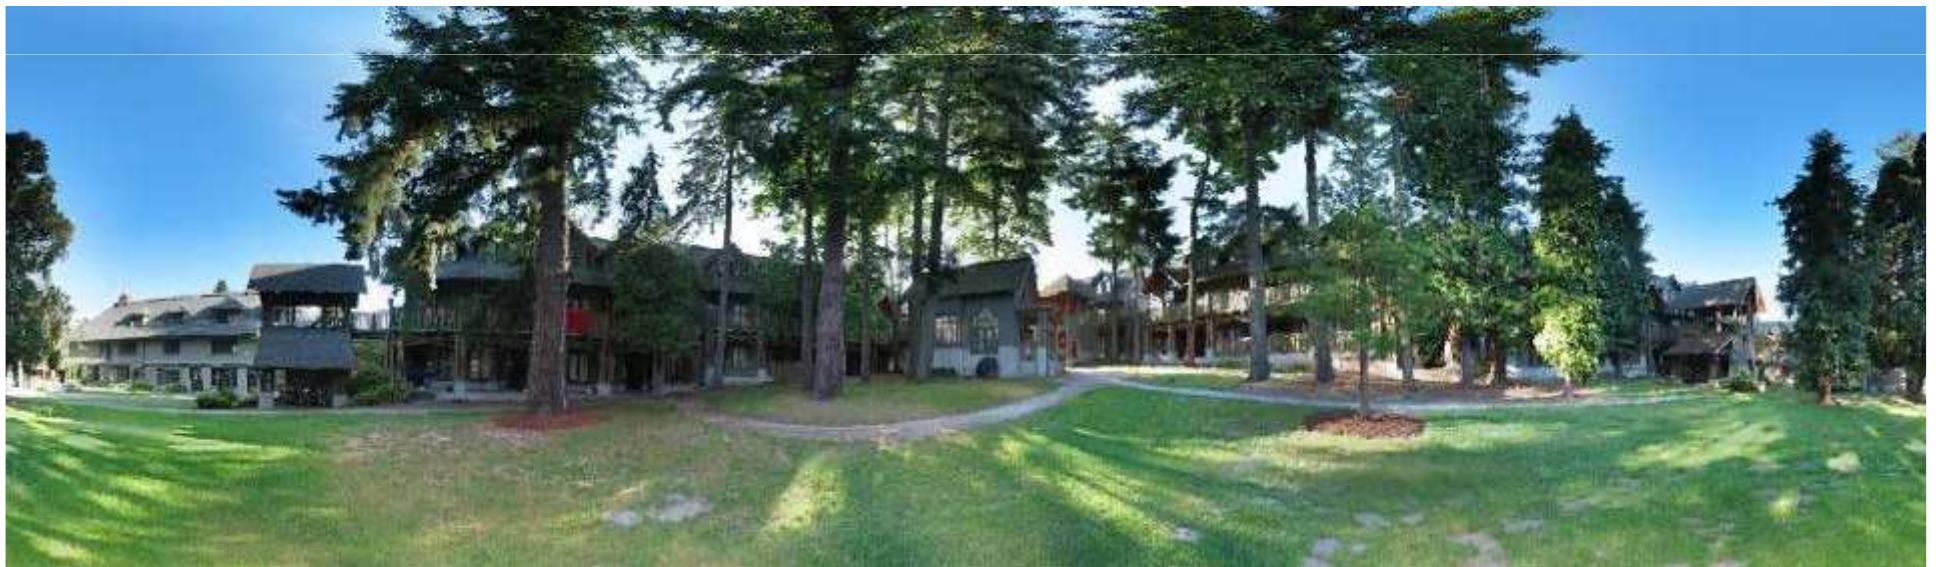
- Compact Camera FOV =  $50 \times 35^\circ$
- Human FOV =  $200 \times 135^\circ$



# Why Mosaic?

Are you getting the whole picture?

- Compact Camera FOV =  $50 \times 35^\circ$
- Human FOV =  $200 \times 135^\circ$
- Panoramic Mosaic =  $360 \times 180^\circ$



Slide from Brown & Lowe

# Mosaics: stitching images together



# Naïve Stitching



left on top



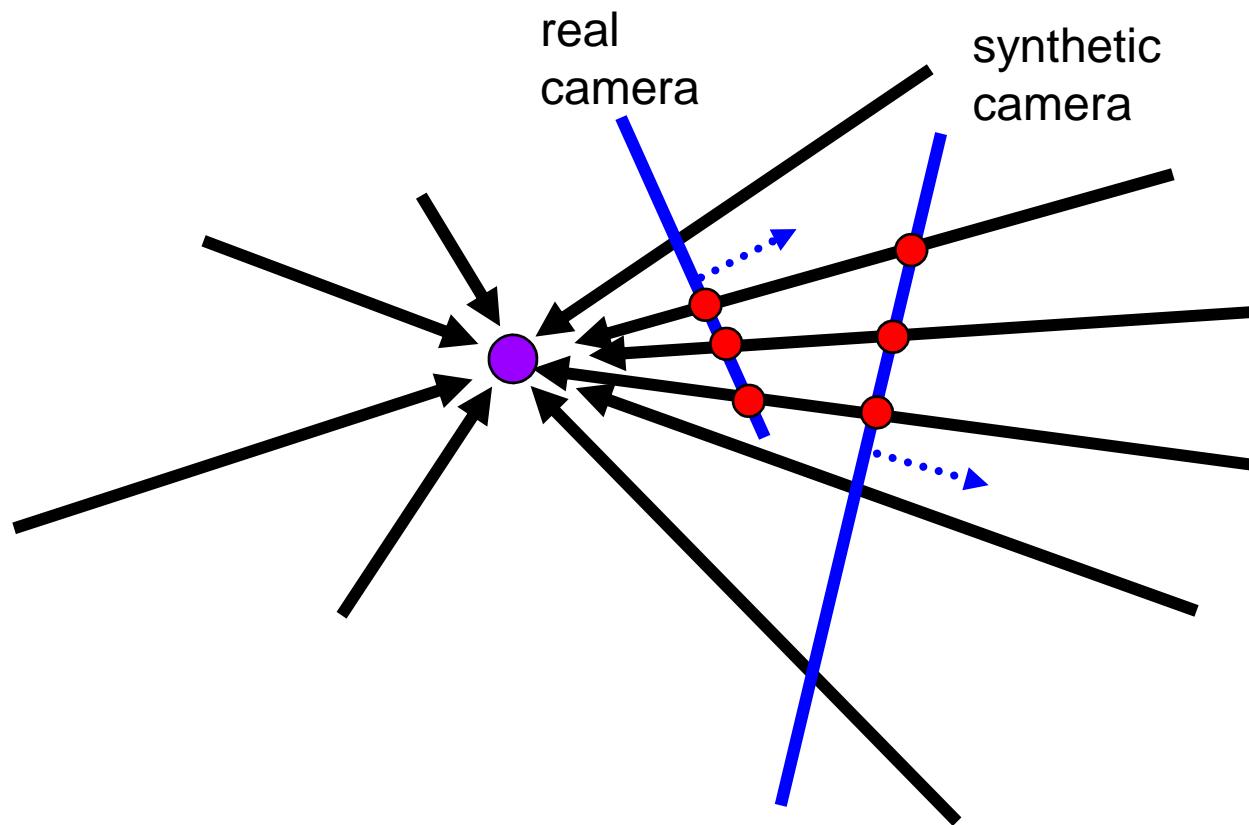
right on top



Translations are not enough to align the images

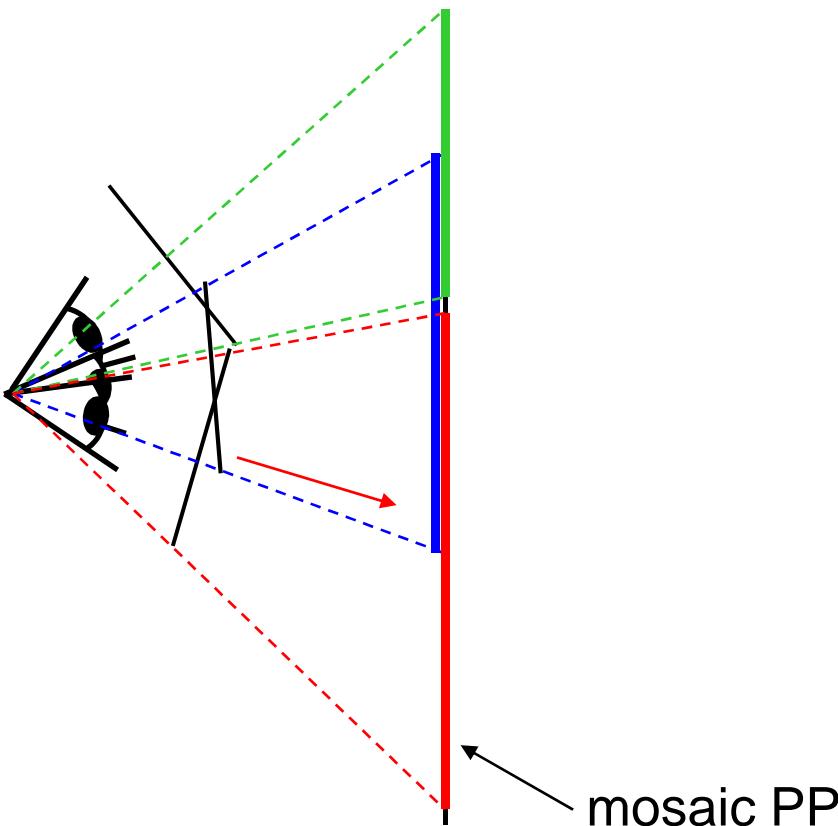


# A pencil of rays contains all views



Can generate any synthetic camera view  
as long as it has **the same center of projection!**

# Image reprojection



The mosaic has a natural interpretation in 3D

- The images are reprojected onto a common plane
- The mosaic is formed on this plane
- Mosaic is a *synthetic wide-angle camera*

# How to do it?

## Basic Procedure

- Take a sequence of images from the same position
  - Rotate the camera about its optical center
- Compute transformation between second image and first
- Transform the second image to overlap with the first
- Blend the two together to create a mosaic
- If there are more images, repeat

...but **wait**, why should this work at all?

- What about the 3D geometry of the scene?
- Why aren't we using it?

# Image reprojection

## Basic question

- How to relate two images from the same camera center?
  - how to map a pixel from PP1 to PP2

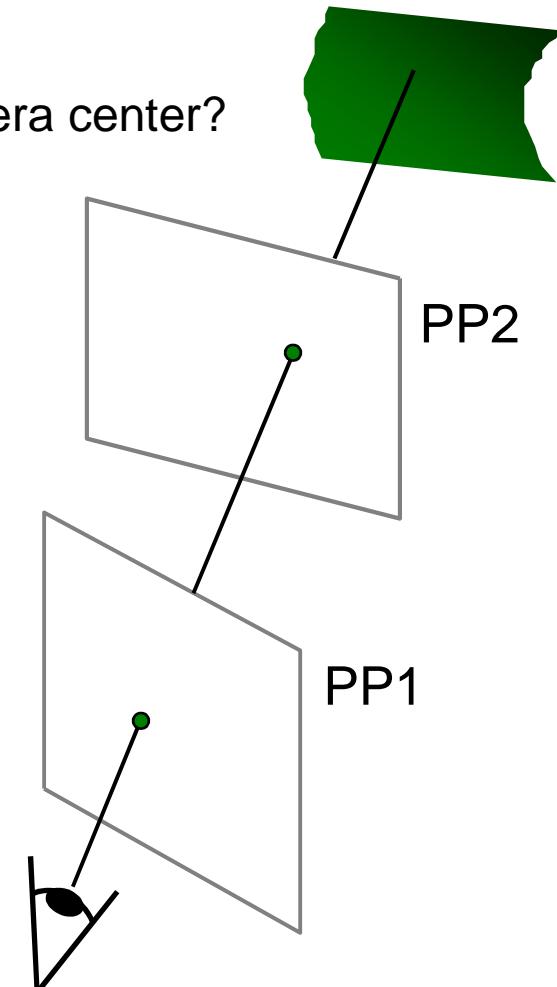
## Answer

- Cast a ray through each pixel in PP1
- Draw the pixel where that ray intersects PP2

But don't we need to know the geometry of the two planes in respect to the eye?

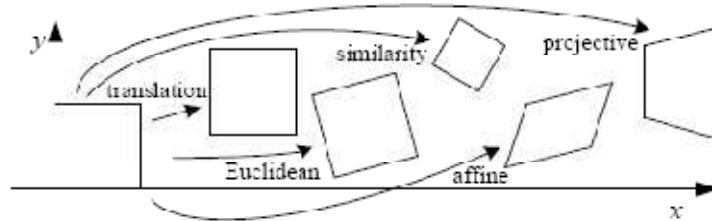
Observation:

Rather than thinking of this as a 3D reprojection, think of it as a 2D **image warp** from one image to another



# Back to Image Warping

Which t-form is the right one for warping PP1 into PP2?  
e.g. translation, Euclidean, affine, projective

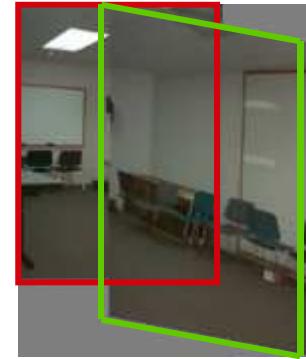


Translation



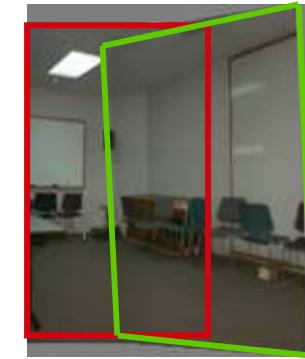
2 unknowns

Affine



6 unknowns

Perspective



8 unknowns

# Homography

A: Projective – mapping between any two PPs with the same center of projection

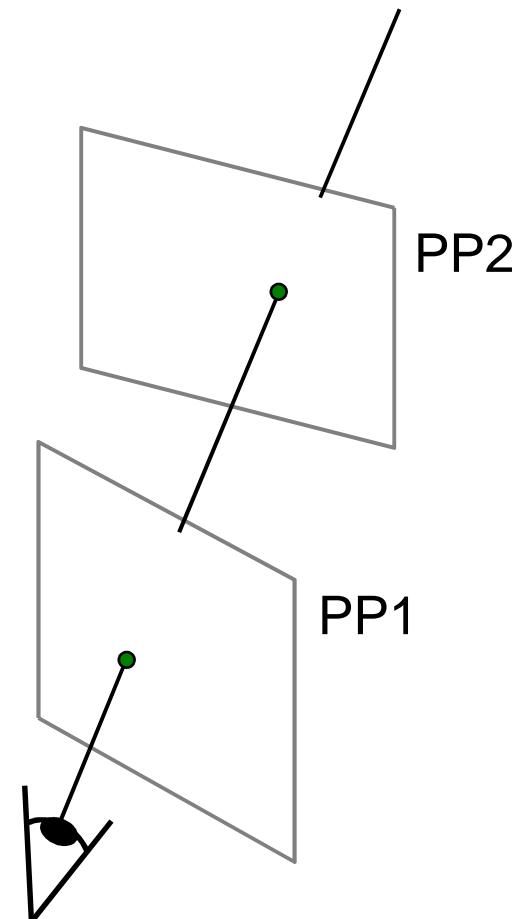
- rectangle should map to arbitrary quadrilateral
- parallel lines aren't
- but must preserve straight lines
- same as: project, rotate, reproject

called Homography

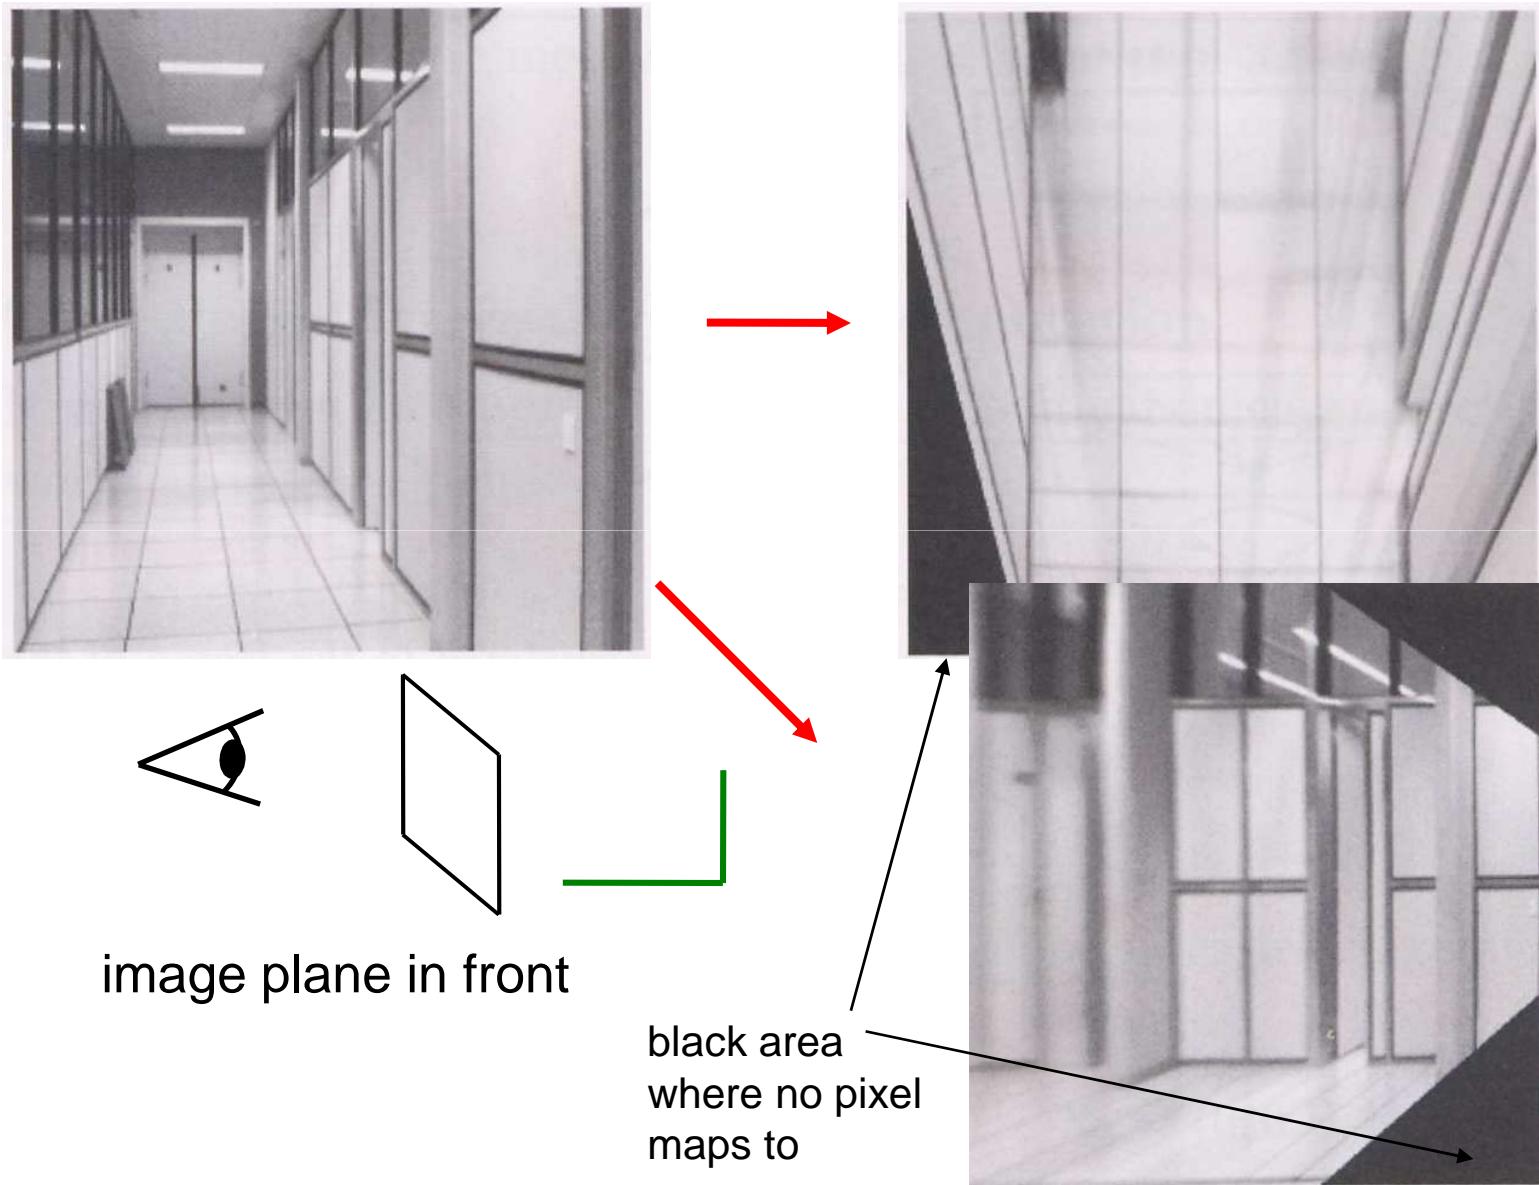
$$\begin{pmatrix} wx' \\ wy' \\ w \\ p' \end{pmatrix} = \begin{pmatrix} a & b & c \\ d & e & f \\ g & h & i \\ \mathbf{H} \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \\ p \end{pmatrix}$$

To apply a homography  $\mathbf{H}$

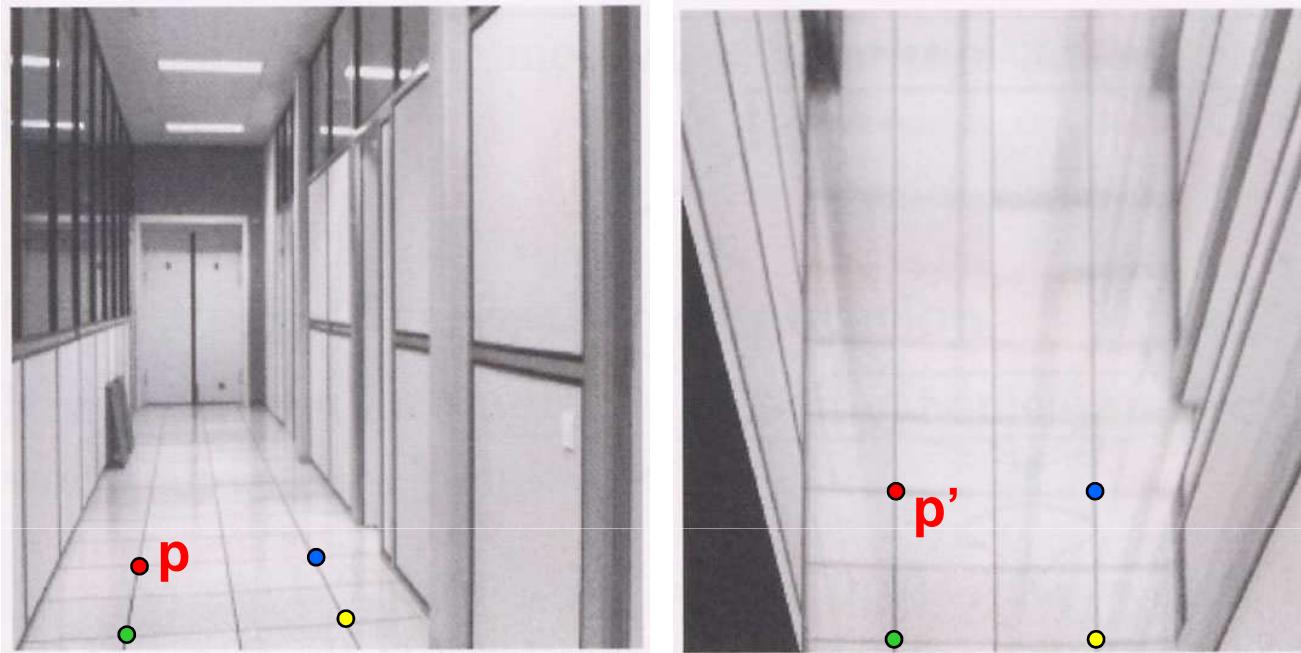
- Compute  $\mathbf{p}' = \mathbf{H}\mathbf{p}$  (regular matrix multiply)
- Convert  $\mathbf{p}'$  from homogeneous to image coordinates



# Image warping with homographies



# Image rectification



To un warp (rectify) an image

- Find the homography  $\mathbf{H}$  given a set of  $\mathbf{p}$  and  $\mathbf{p}'$  pairs
- How many correspondences are needed?
- Tricky to write  $\mathbf{H}$  analytically, but we can solve for it!
  - Find such  $\mathbf{H}$  that “best” transforms points  $\mathbf{p}$  into  $\mathbf{p}'$
  - Use least-squares!

# Least Squares Example

Say we have a set of data points  $(X_1, X_1')$ ,  $(X_2, X_2')$ ,  $(X_3, X_3')$ , etc. (e.g. person's height vs. weight)

We want a nice compact formula (a line) to predict  $X'$ 's from  $X$ 's:  $Xa + b = X'$

We want to find  $a$  and  $b$

How many  $(X, X')$  pairs do we need?

$$X_1 a + b = X_1'$$

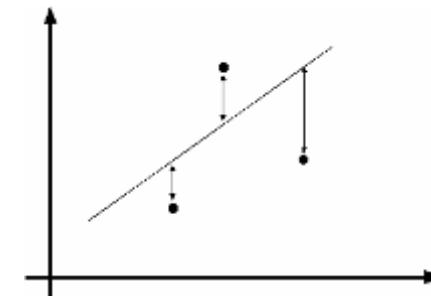
$$X_2 a + b = X_2'$$

$$\begin{bmatrix} X_1 & 1 \\ X_2 & 1 \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} X_1' \\ X_2' \end{bmatrix} \quad Ax=B$$

What if the data is noisy?

$$\begin{bmatrix} X_1 & 1 \\ X_2 & 1 \\ X_3 & 1 \\ \dots & \dots \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} X_1' \\ X_2' \\ X_3' \\ \dots \end{bmatrix}$$

$$\min \|Ax - B\|^2$$



overconstrained

# Solving for homographies

$$\begin{aligned} \mathbf{p}' &= \mathbf{H}\mathbf{p} \\ \begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} &= \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \end{aligned}$$

Can set scale factor  $i=1$ . So, there are 8 unknowns.

Set up a system of linear equations:

$$\mathbf{A}\mathbf{h} = \mathbf{b}$$

where vector of unknowns  $\mathbf{h} = [a, b, c, d, e, f, g, h]^T$

Need at least 8 eqs, but the more the better...

Solve for  $\mathbf{h}$ . If overconstrained, solve using least-squares:

$$\min \| \mathbf{A}\mathbf{h} - \mathbf{b} \|^2$$

Can be done in Matlab using “\” command

- see “help lmdivide”

# Fun with homographies

Original image



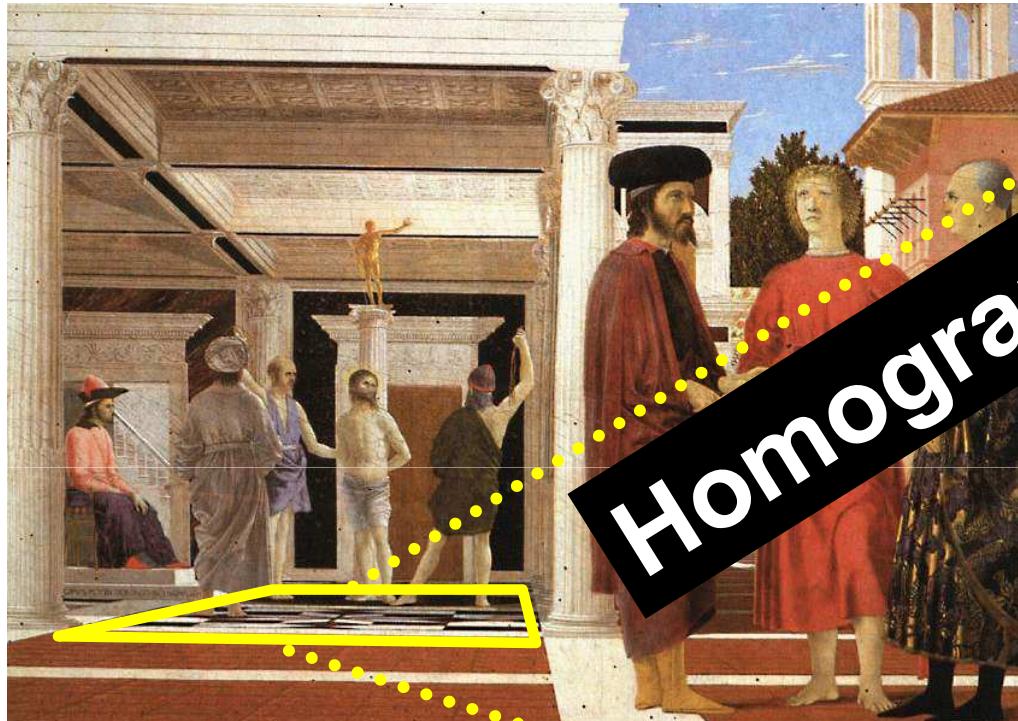
St.Petersburg  
photo by A. Tikhonov

Virtual camera rotations



# Analysing patterns and shapes

What is the shape of the b/w floor pattern?



The floor (enlarged)



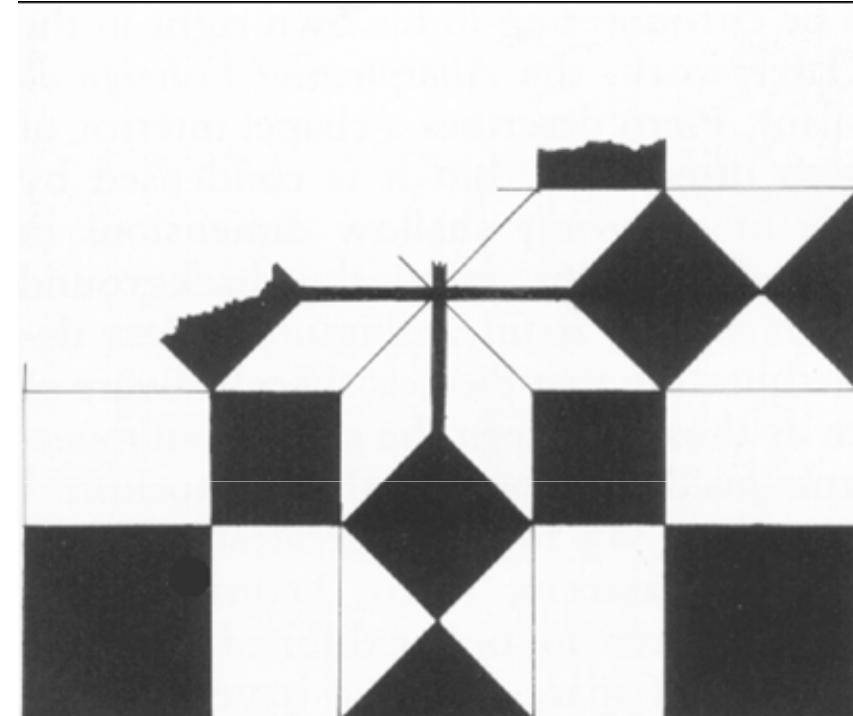
Automatically  
rectified floor

# Analysing patterns and shapes

Automatic rectification



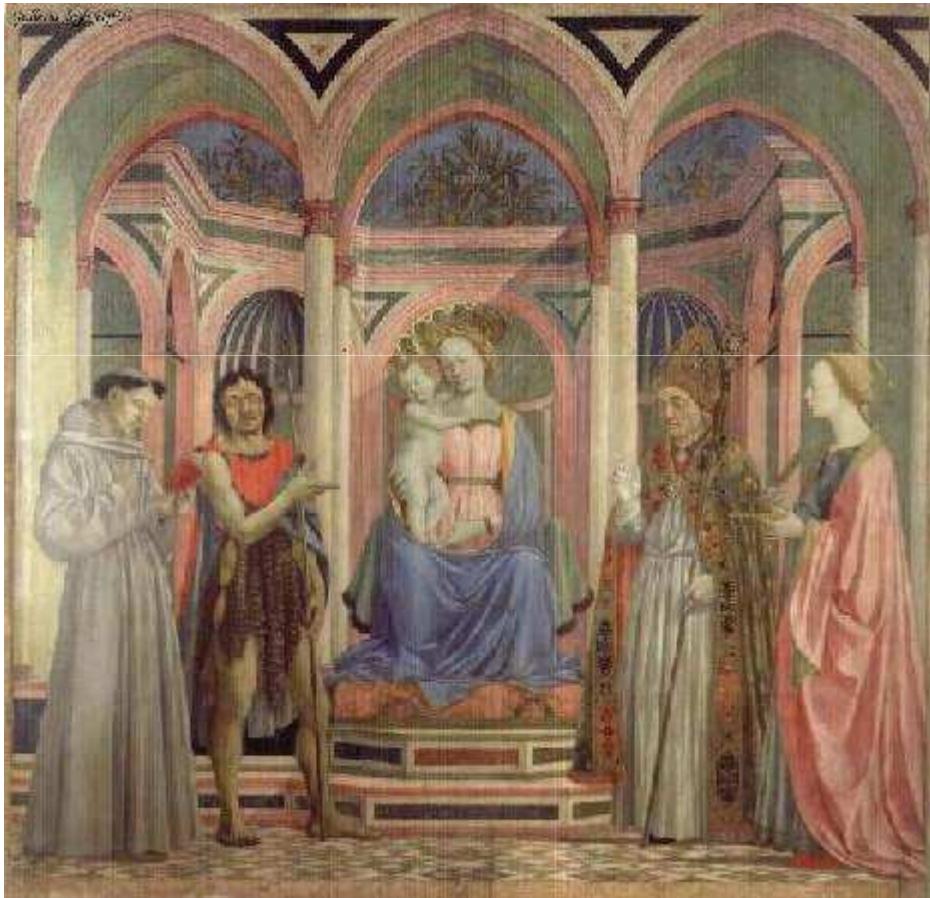
Slide from Criminisi



From Martin Kemp *The Science of Art*  
*(manual reconstruction)*

2 patterns have been discovered !

# Analysing patterns and shapes



What is the (complicated) shape of the floor pattern?



Automatically rectified floor

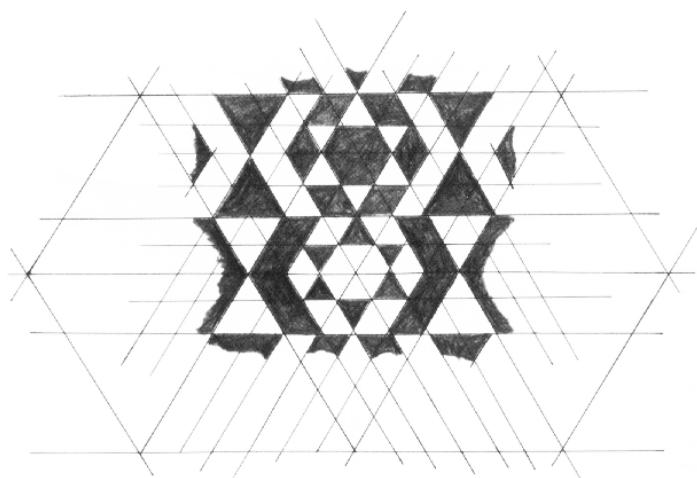
***St. Lucy Altarpiece, D. Veneziano***

Slide from Criminisi

# Analysing patterns and shapes



**Automatic  
rectification**



**From Martin Kemp, *The Science of Art*  
(manual reconstruction)**

# Julian Beever: Manual Homographies

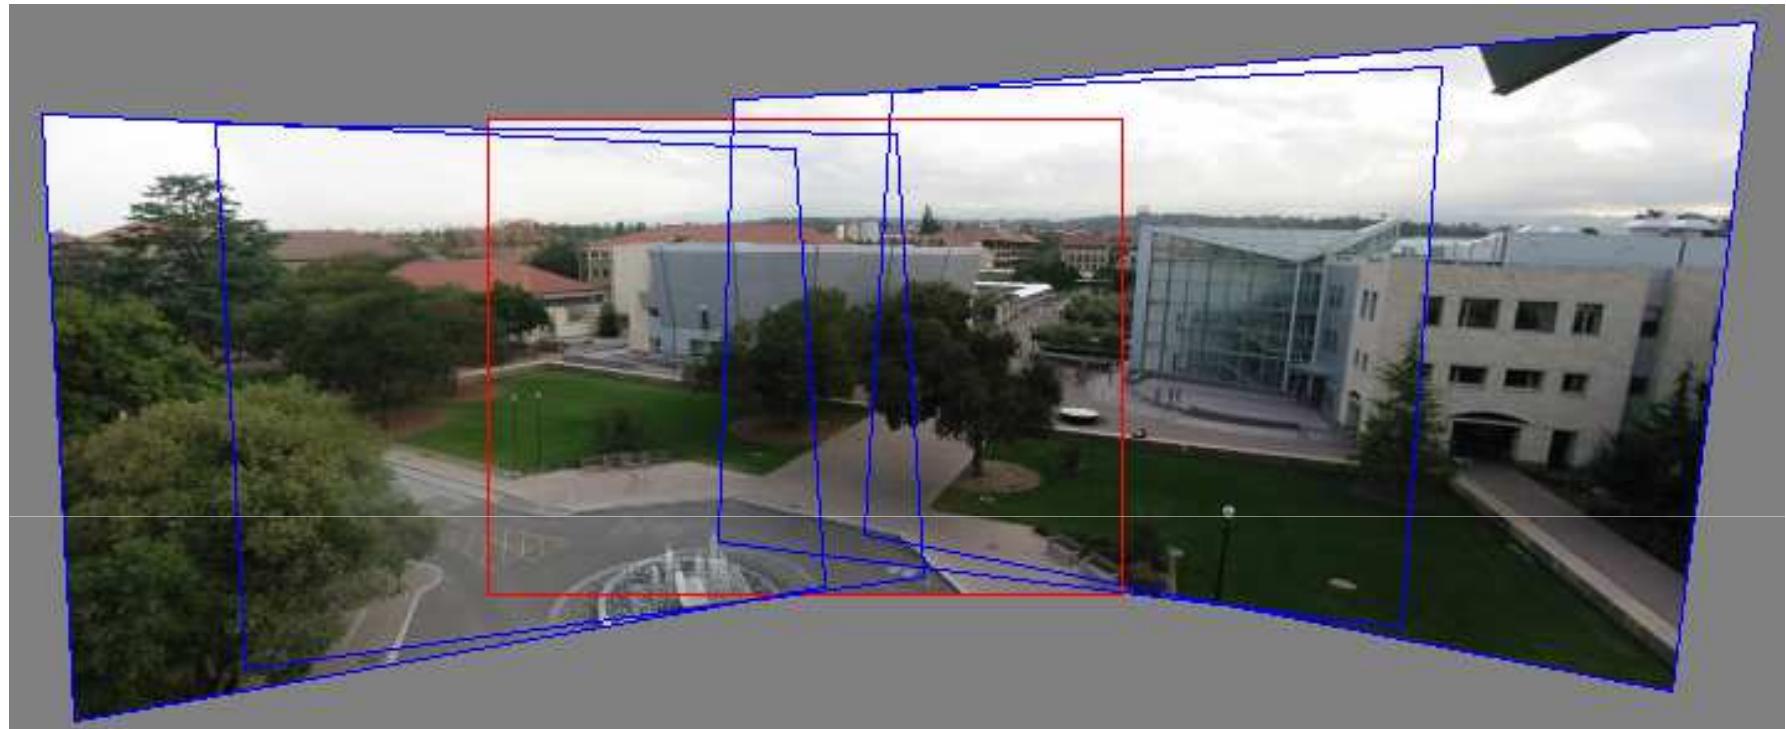


<http://users.skynet.be/J.Beever/pave.htm>

# Holbein, *The Ambassadors*



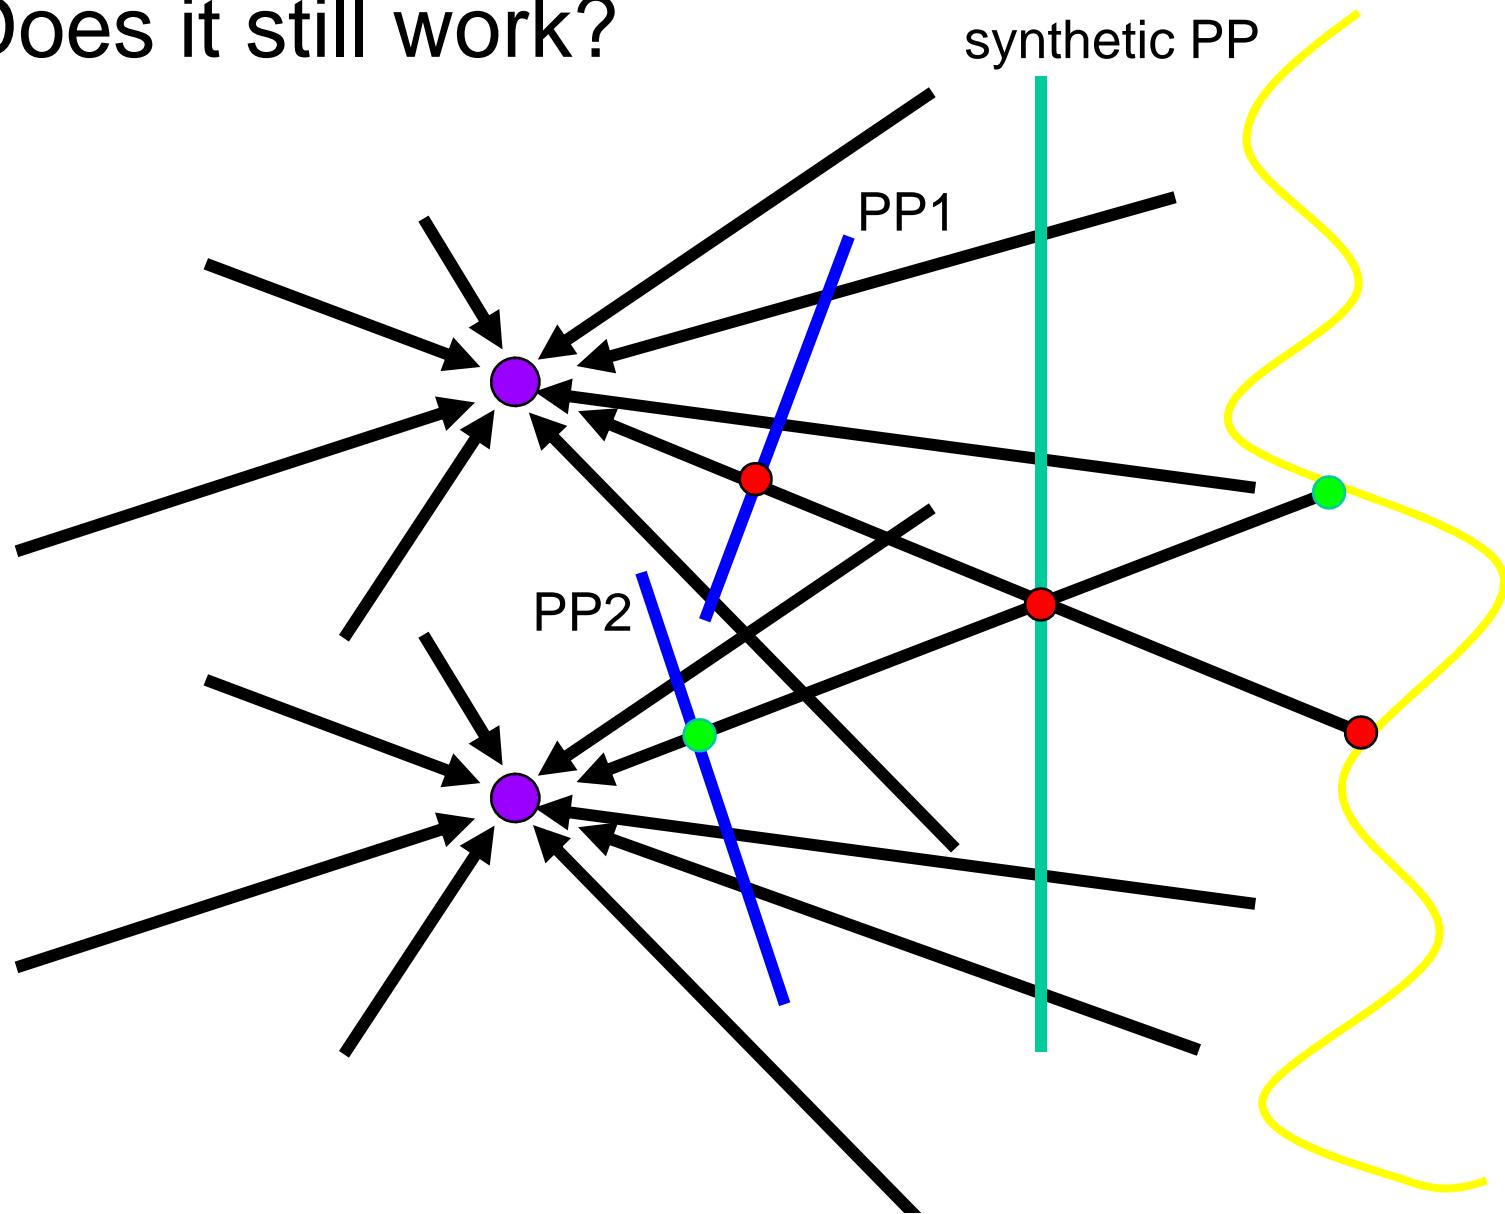
# Panoramas



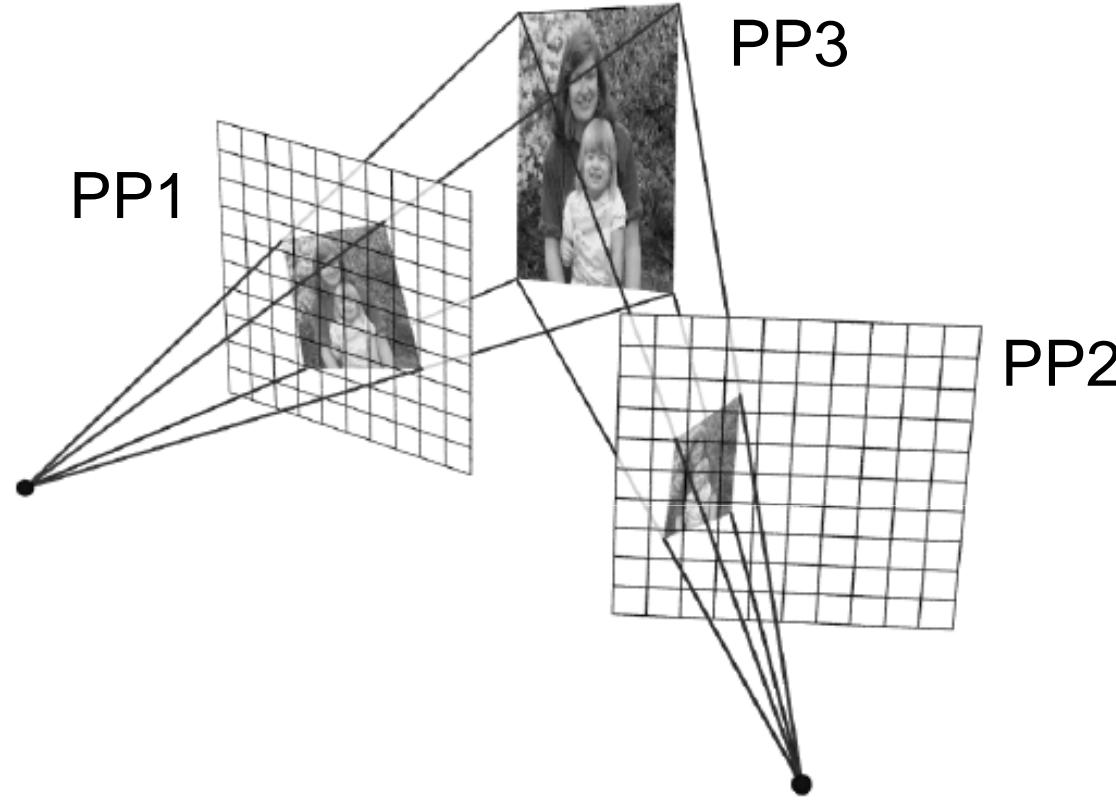
1. Pick one image (red)
2. Warp the other images towards it (usually, one by one)
3. Blend

# changing camera center

Does it still work?



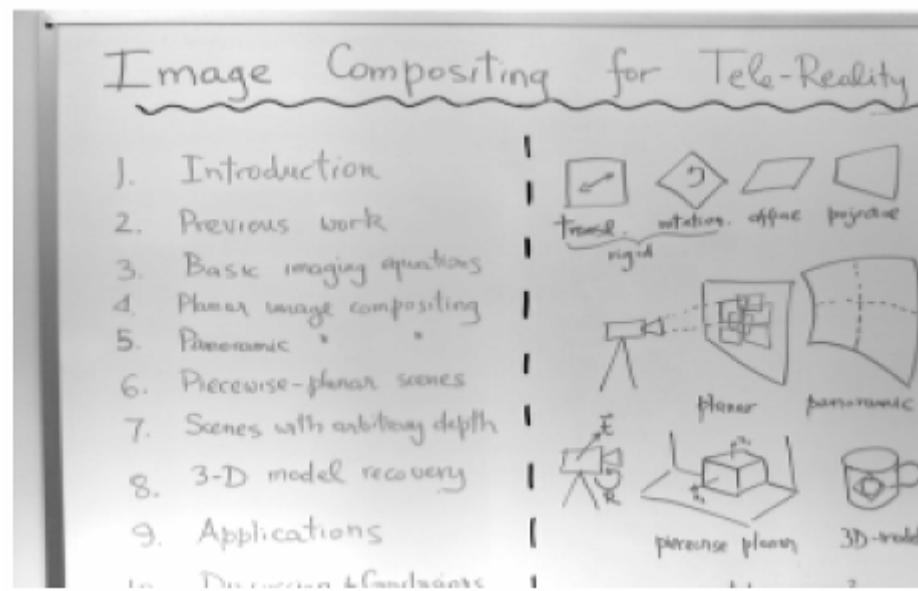
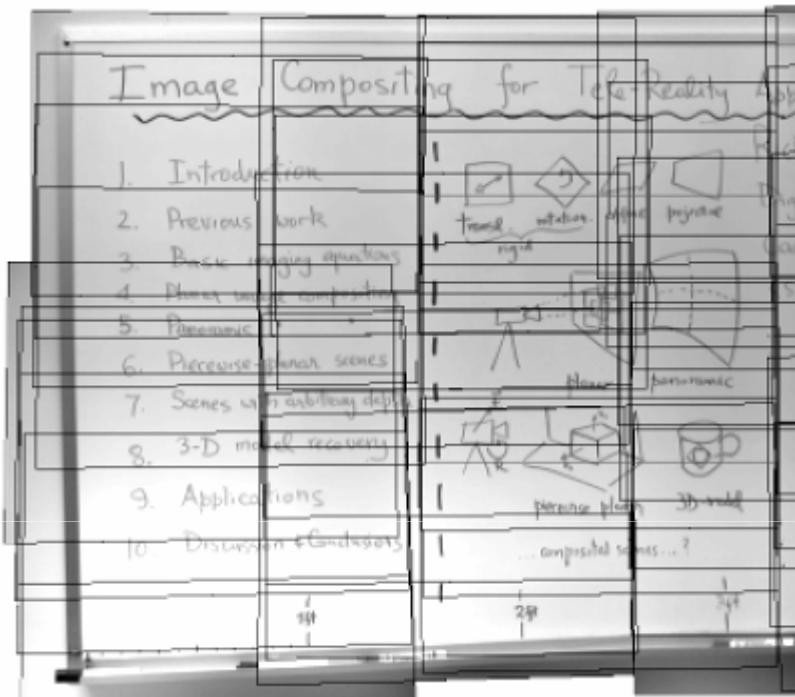
# Planar scene (or far away)



PP3 is a projection plane of both centers of projection,  
so we are OK!

This is how big aerial photographs are made

# Planar mosaic



# Some examples at CMU



Ben Hollis, 2004



Ben Hollis, 2004

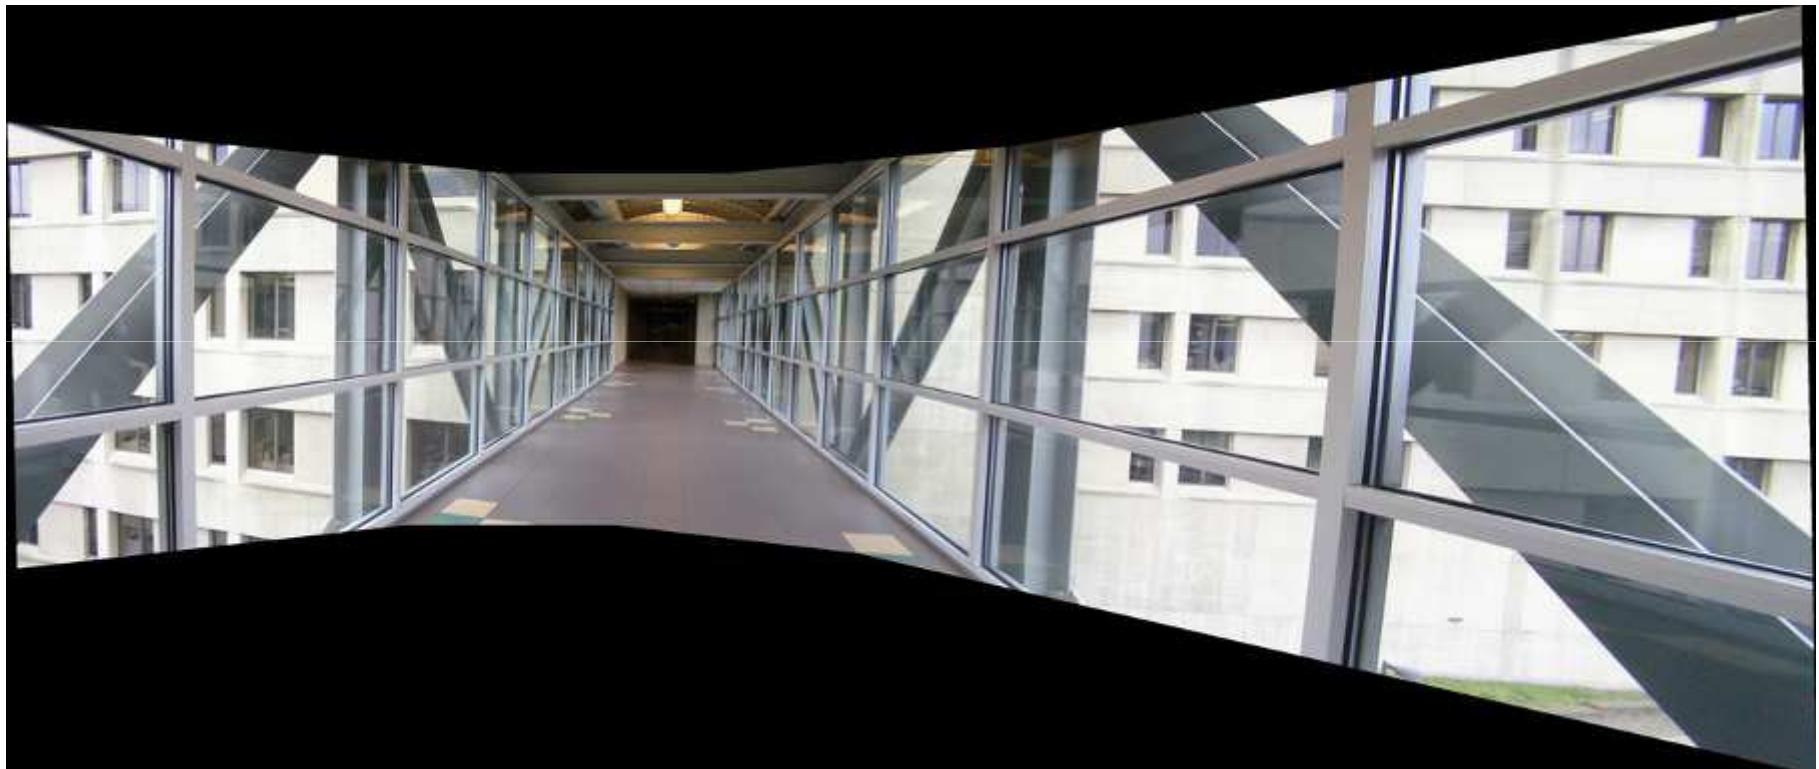


Matt Pucevich , 2004



Eunjeong Ryu (E.J), 2004

# Go Explore!



# Breaking out of 2D

...now we are ready to break out of 2D



And enter the real world!



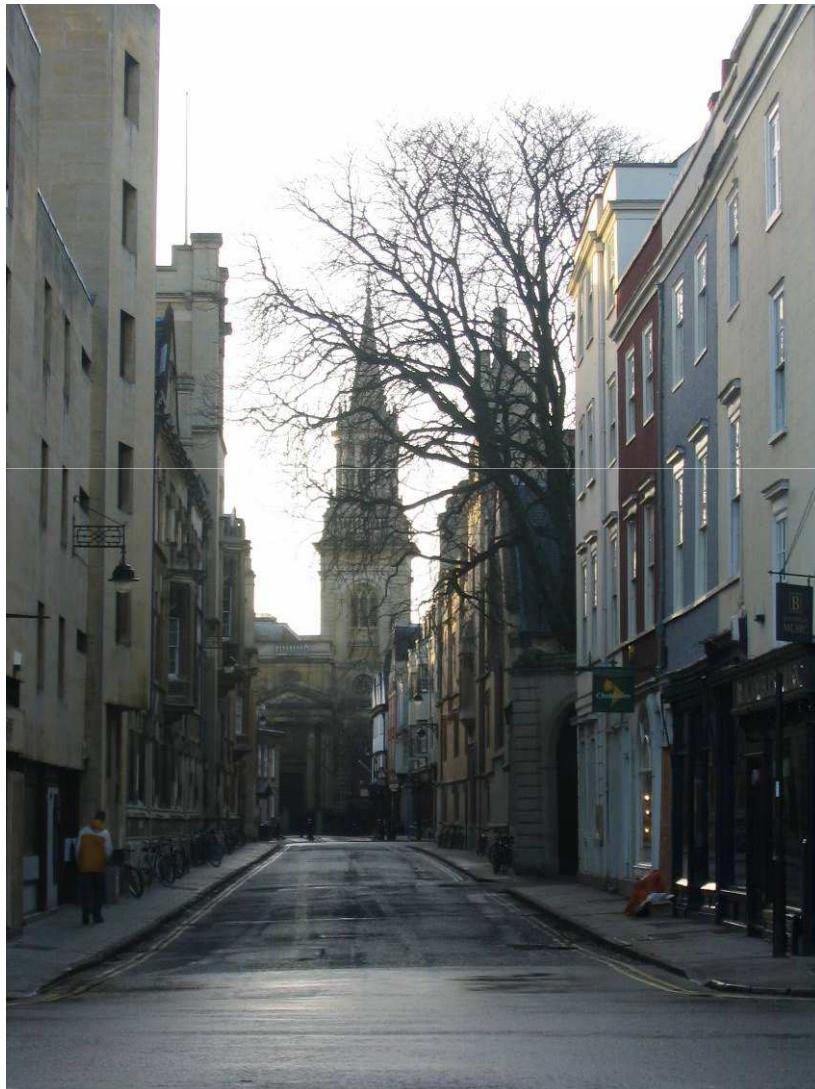
# on to 3D...

Enough of images!

We want real 3D scene  
walk-throughs:

- Camera rotation
- Camera translation

Can we do it from a  
single photograph?



# Camera rotations with homographies

Original image



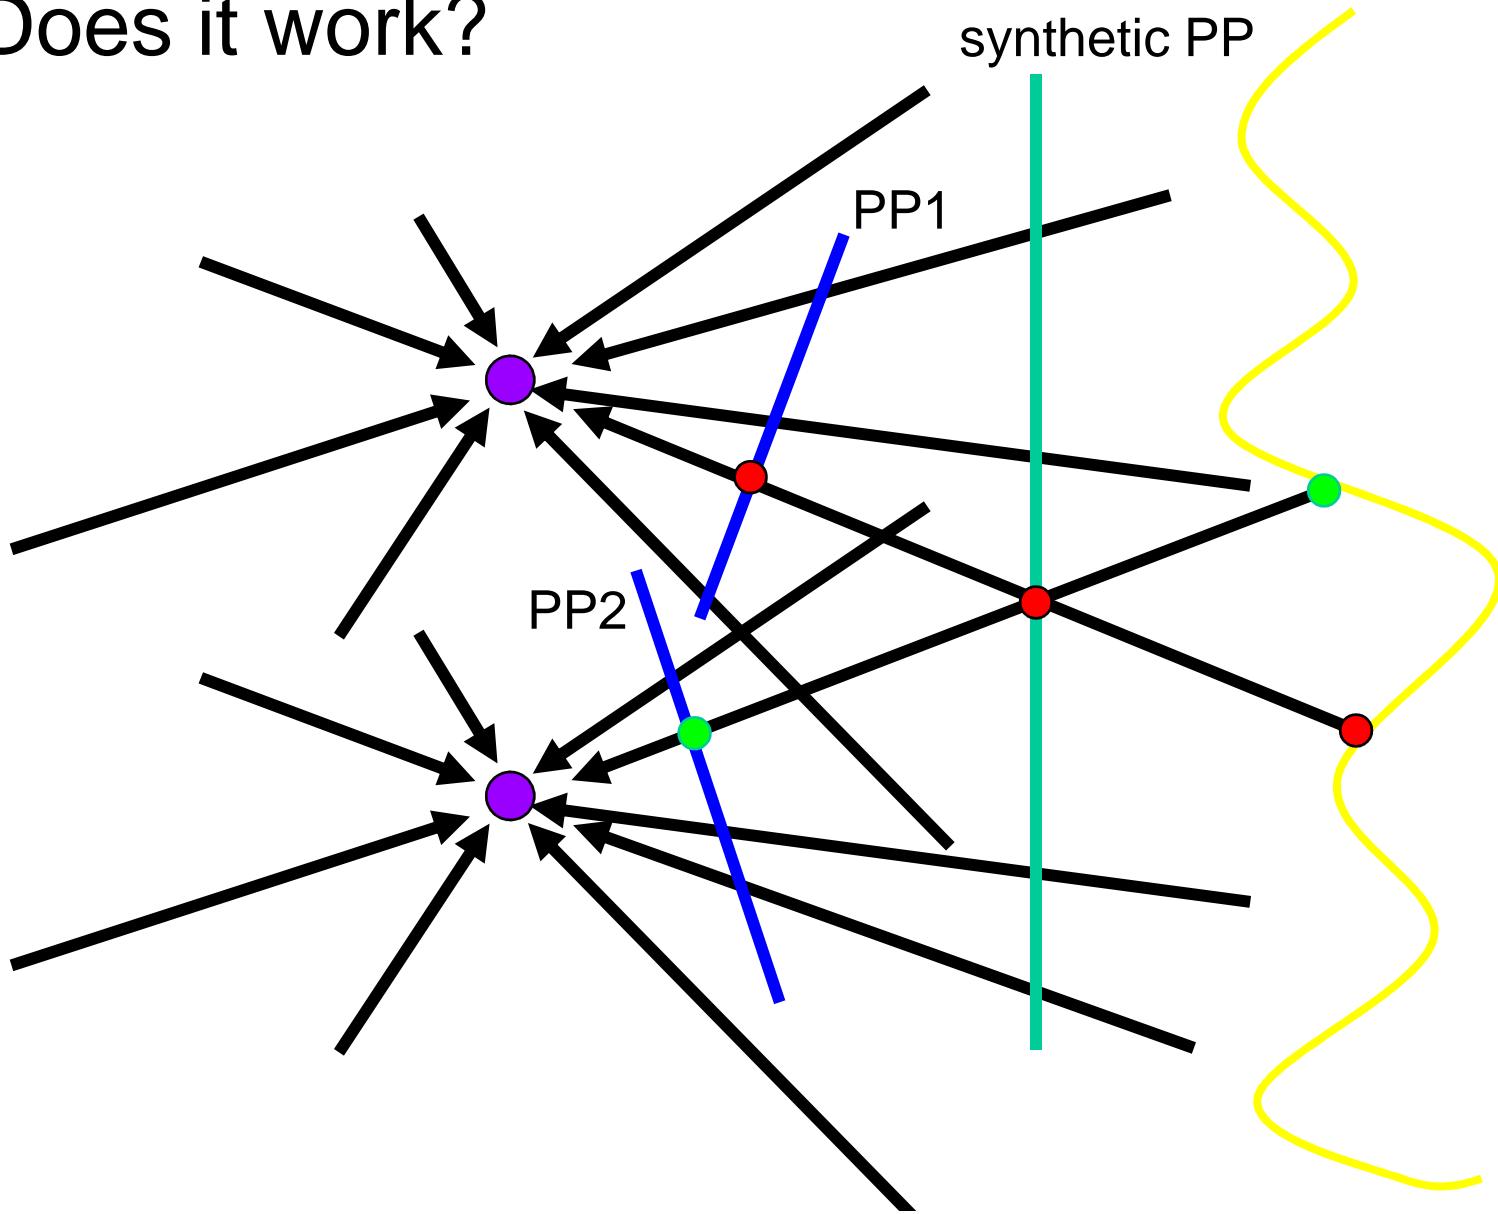
St.Petersburg  
photo by A. Tikhonov

Virtual camera rotations

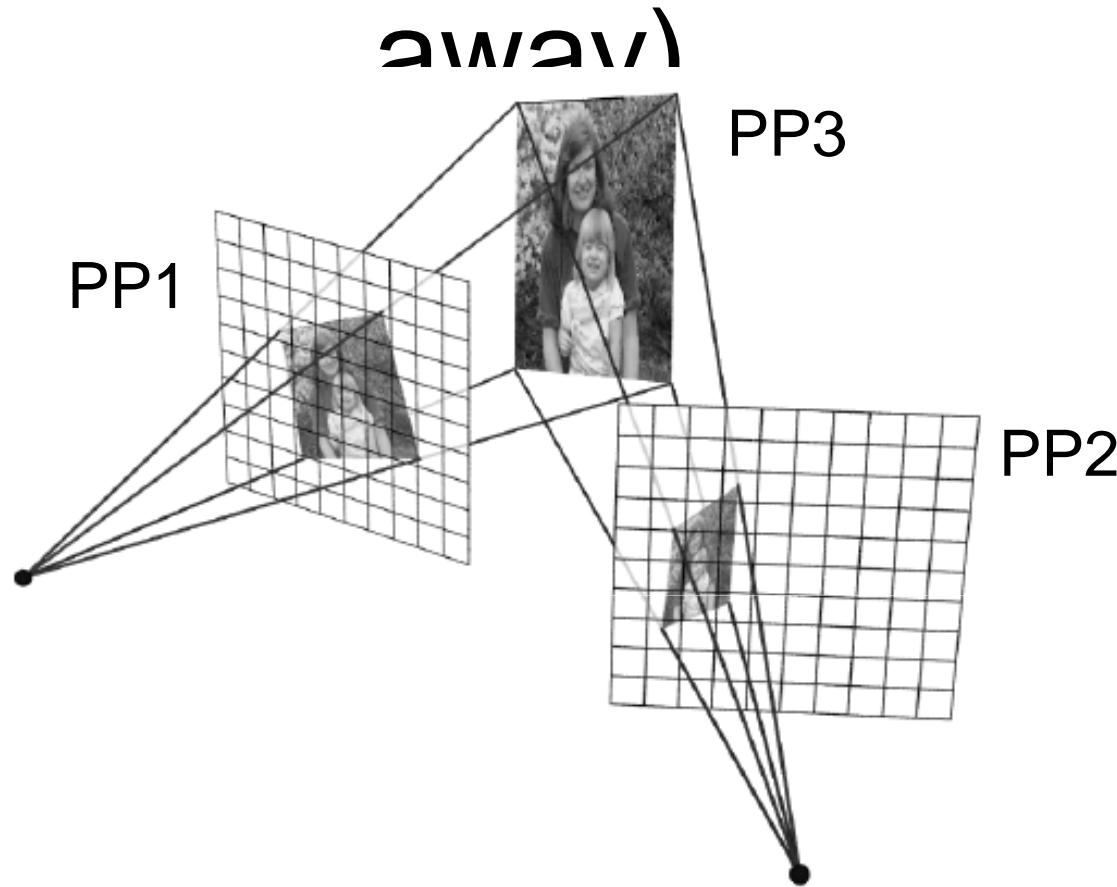


# Camera translation

Does it work?



# Yes, with planar scene (or far away)

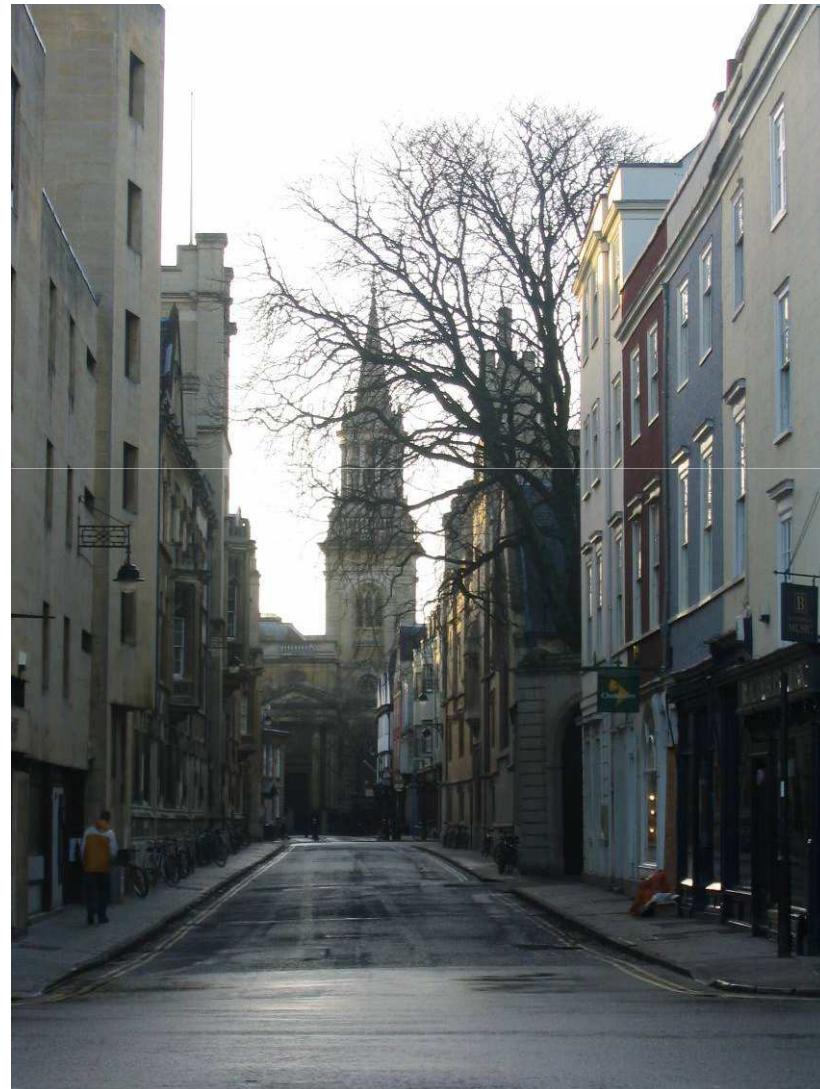


PP3 is a projection plane of both centers of projection, so we are OK!

# So, what can we do here?

Model the scene as a set of planes!

Now, just need to find the orientations of these planes.



# Some preliminaries: projective geometry



[Ames Room](#)

# Silly Euclid: Trix are for kids!



Parallel lines???

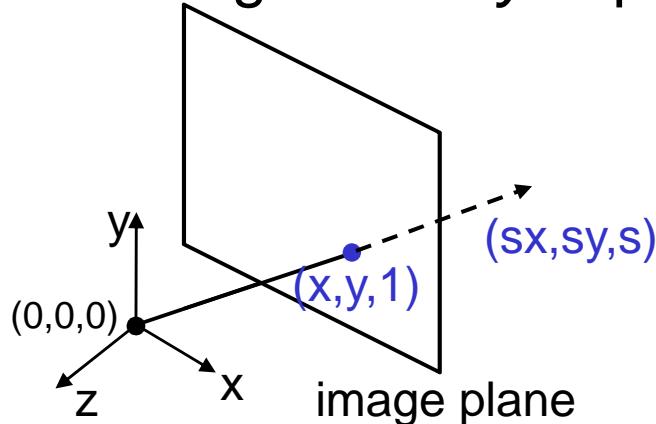
# The projective plane

Why do we need homogeneous coordinates?

- represent points at infinity, homographies, perspective projection, multi-view relationships

What is the geometric intuition?

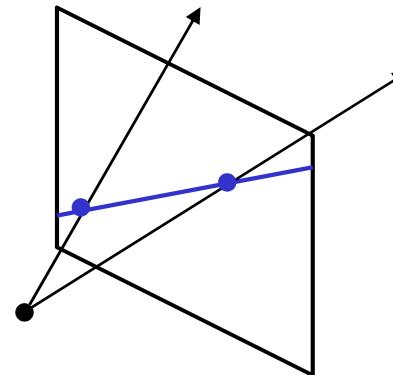
- a point in the image is a *ray* in projective space



- Each *point*  $(x,y)$  on the plane is represented by a *ray*  $(sx, sy, s)$ 
  - all points on the ray are equivalent:  $(x, y, 1) \cong (sx, sy, s)$

# Projective lines

What does a line in the image correspond to in projective space?



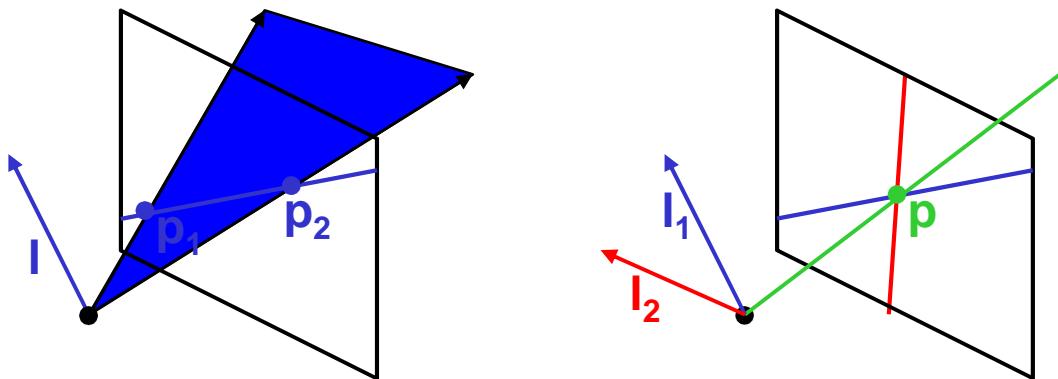
- A line is a *plane* of rays through origin
  - all rays  $(x,y,z)$  satisfying:  $ax + by + cz = 0$

$$0 = \begin{pmatrix} a & b & c \\ l & & \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ p \end{pmatrix}$$

- A line is also represented as a homogeneous 3-vector  $\mathbf{l}$

# Point and line duality

- A line  $\mathbf{l}$  is an homogeneous 3-vector
- It is  $\perp$  to every point (ray) on the line  $\mathbf{l}\mathbf{p} = 0$



What is the line  $\mathbf{l}$  spanned by the rays  $\mathbf{p}_1$  and  $\mathbf{p}_2$ ?

- $\mathbf{l}$  is  $\perp$  to  $\mathbf{p}_1$  and  $\mathbf{p}_2 \Rightarrow \mathbf{l} = \mathbf{p}_1 \times \mathbf{p}_2$
- $\mathbf{l}$  is the plane normal vector

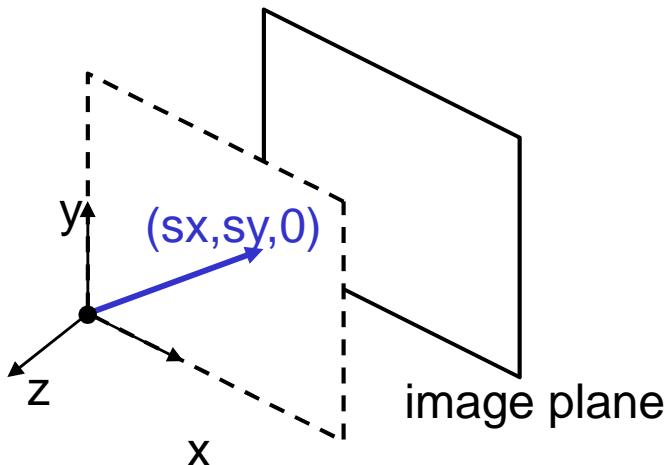
What is the intersection of two lines  $\mathbf{l}_1$  and  $\mathbf{l}_2$ ?

- $\mathbf{p}$  is  $\perp$  to  $\mathbf{p}_1$  and  $\mathbf{p}_2 \Rightarrow \mathbf{l} = \mathbf{p}_1 \times \mathbf{p}_2$

Points and lines are dual in the projective space

- Given any formula, we can switch the meaning of points and lines to get another one.

# Ideal points and lines



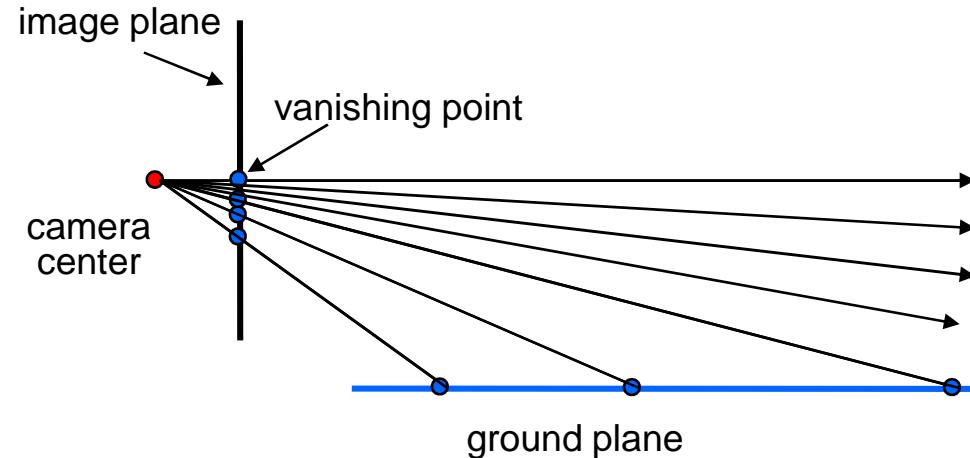
Ideal point (“point at infinity”)

- $p \cong (x, y, 0)$  – parallel to image plane
- It has infinite image coordinates

Ideal line

$l \cong (0, 0, 1)$  – parallel to image plane

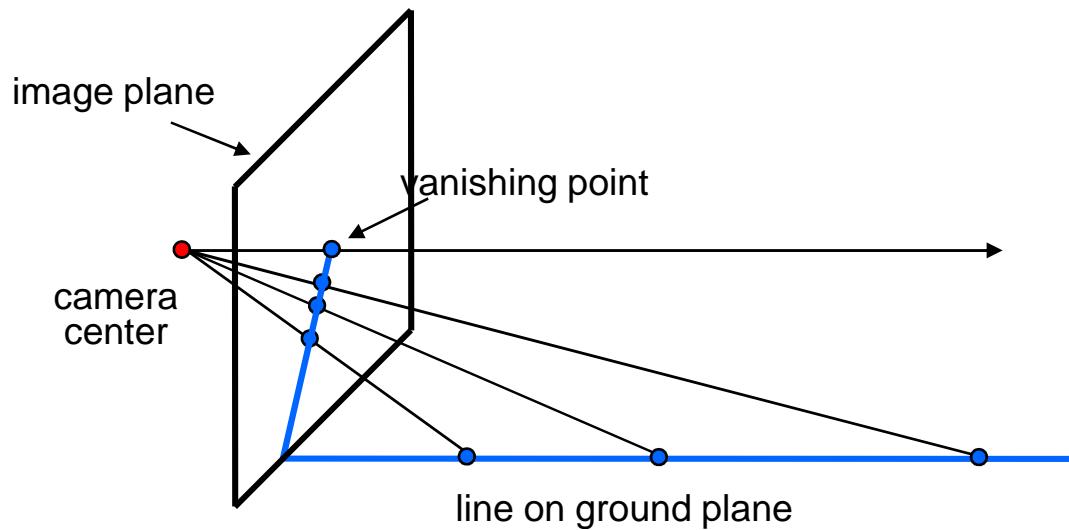
# Vanishing points



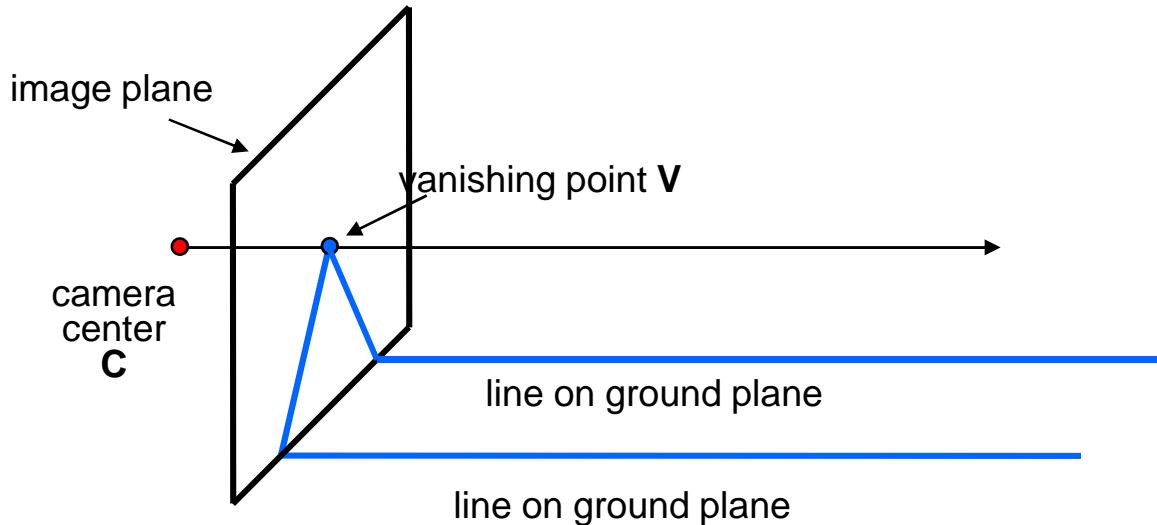
## Vanishing point

- projection of a point at infinity
- Caused by ideal line

# Vanishing points (2D)



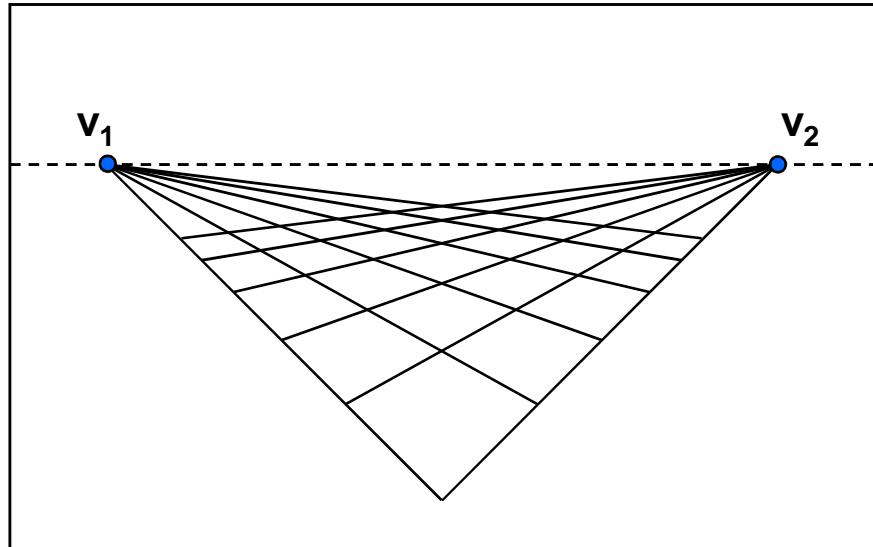
# Vanishing points



## Properties

- Any two parallel lines have the same vanishing point **v**
- The ray from **C** through **v** is parallel to the lines
- An image may have more than one vanishing point

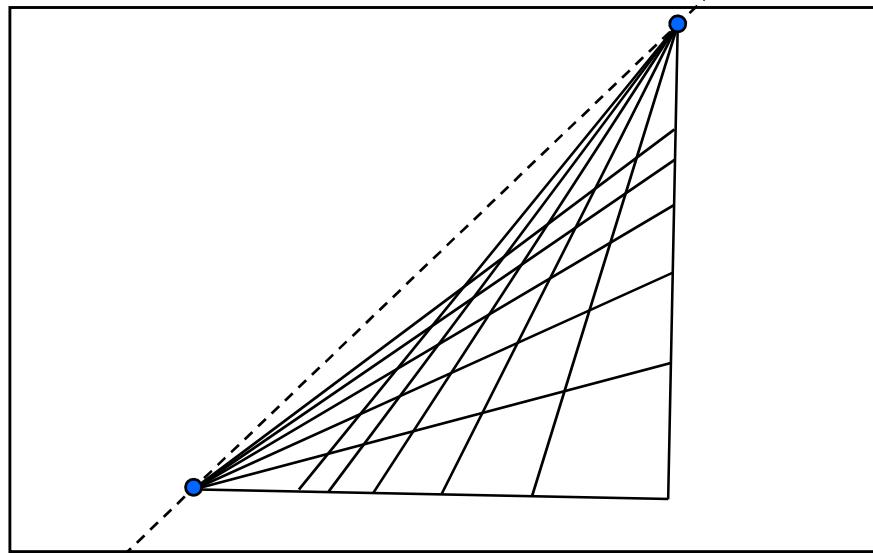
# Vanishing lines



## Multiple Vanishing Points

- Any set of parallel lines on the plane define a vanishing point
- The union of all of these vanishing points is the *horizon line*
  - also called *vanishing line*
- Note that different planes define different vanishing lines

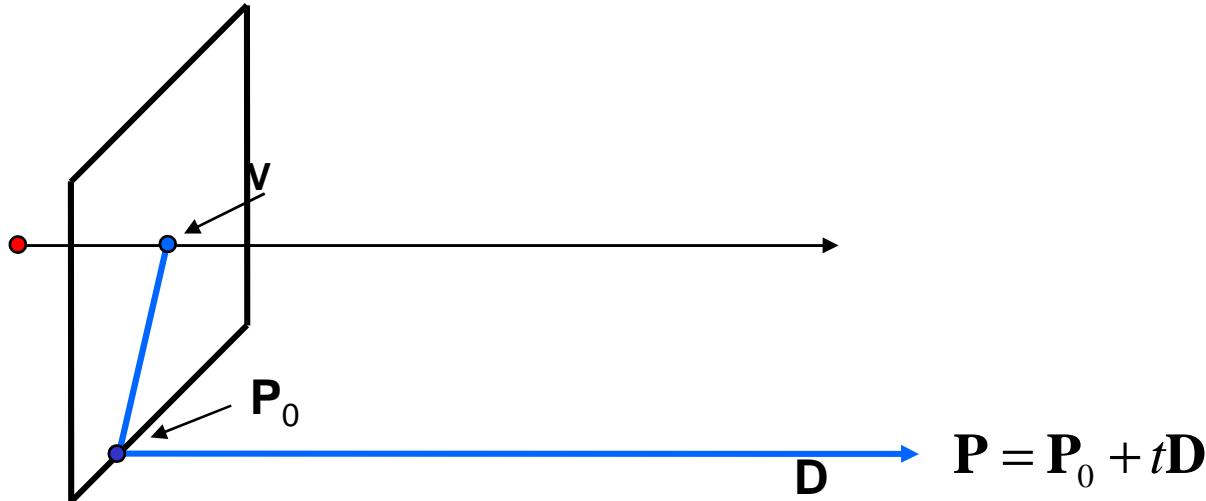
# Vanishing lines



## Multiple Vanishing Points

- Any set of parallel lines on the plane define a vanishing point
- The union of all of these vanishing points is the *horizon line*
  - also called *vanishing line*
- Note that different planes define different vanishing lines

# Computing vanishing points

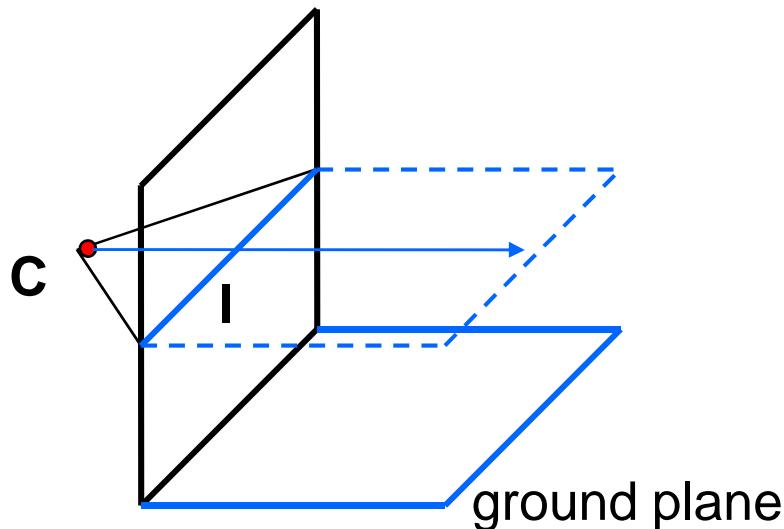


$$\mathbf{P}_t = \begin{bmatrix} P_x + tD_x \\ P_y + tD_y \\ P_z + tD_z \\ 1 \end{bmatrix} \cong \begin{bmatrix} P_x / t + D_x \\ P_y / t + D_y \\ P_z / t + D_z \\ 1/t \end{bmatrix} \quad t \rightarrow \infty \quad \mathbf{P}_\infty \cong \begin{bmatrix} D_x \\ D_y \\ D_z \\ 0 \end{bmatrix}$$

## Properties $\mathbf{v} = \Pi \mathbf{P}_\infty$

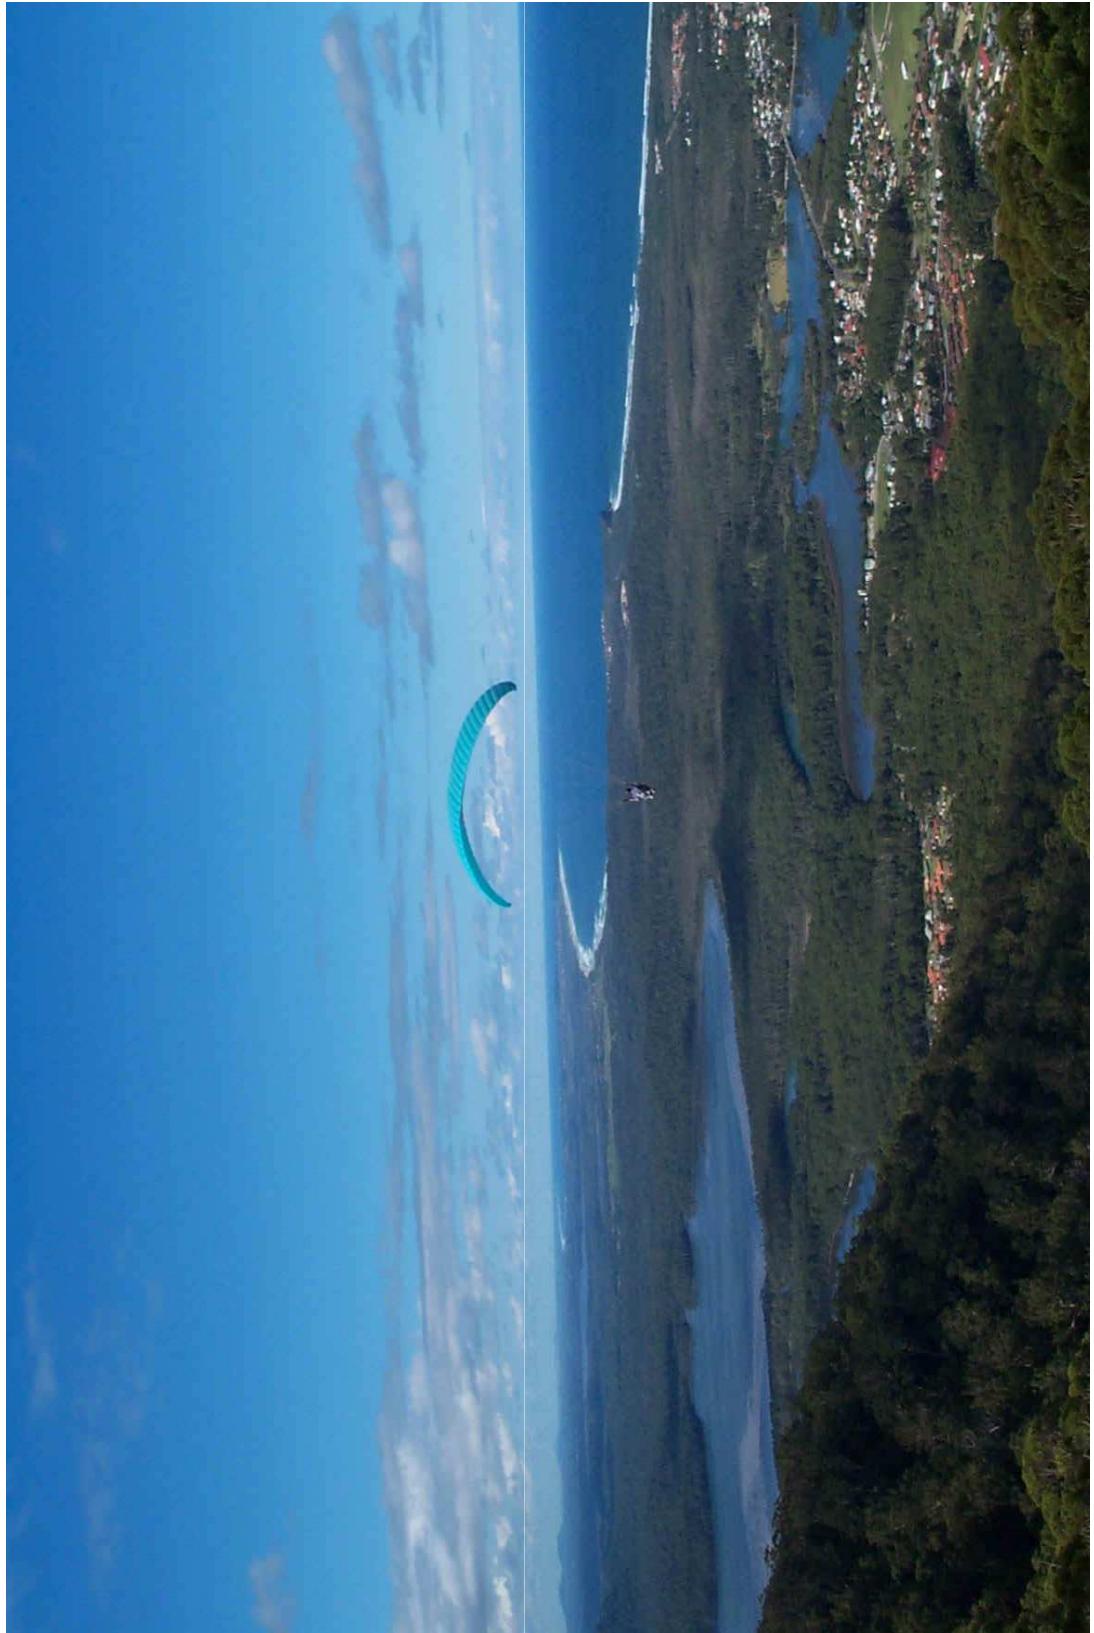
- $\mathbf{P}_\infty$  is a point at *infinity*,  $\mathbf{v}$  is its projection
- They depend only on line *direction*
- Parallel lines  $\mathbf{P}_0 + t\mathbf{D}$ ,  $\mathbf{P}_1 + t\mathbf{D}$  intersect at  $\mathbf{P}_\infty$

# Computing vanishing lines

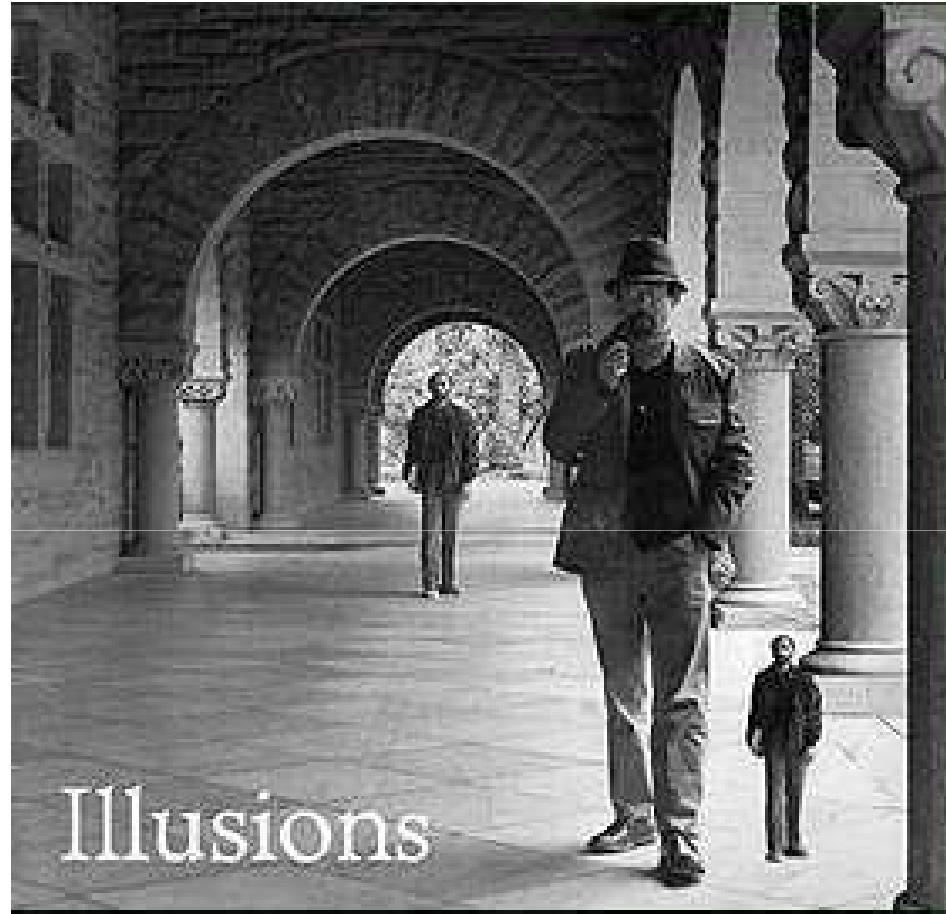
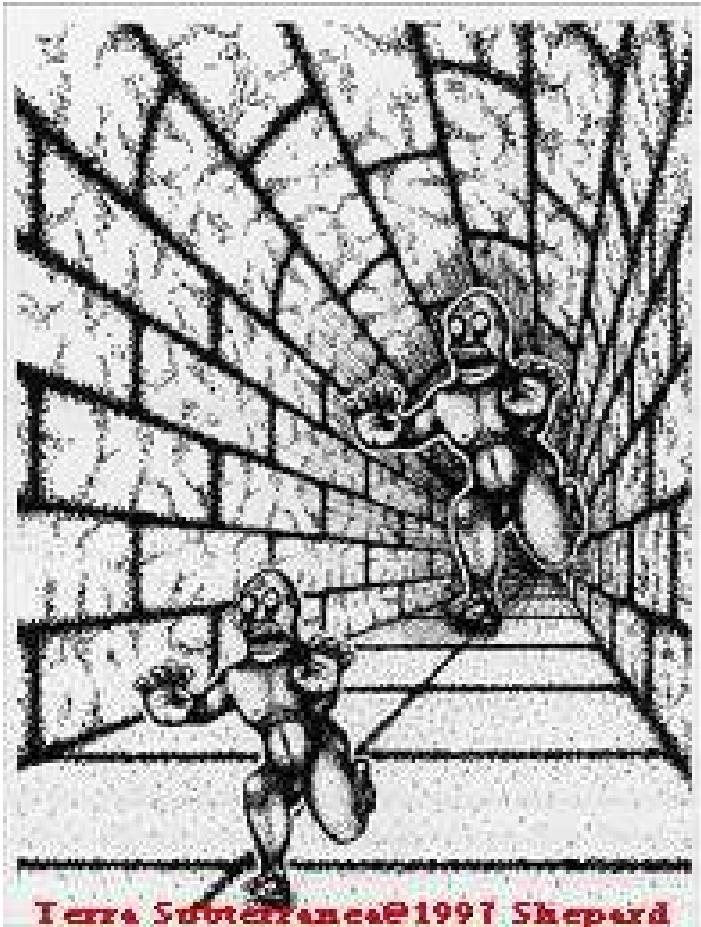


## Properties

- **I** is intersection of horizontal plane through **C** with image plane
- Compute **I** from two sets of parallel lines on ground plane
- All points at same height as **C** project to **I**
  - points higher than **C** project above **I**
- Provides way of comparing height of objects in the scene

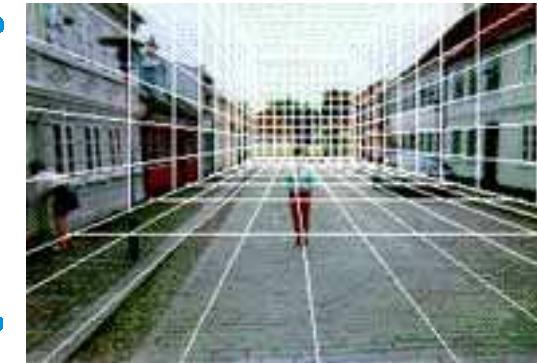
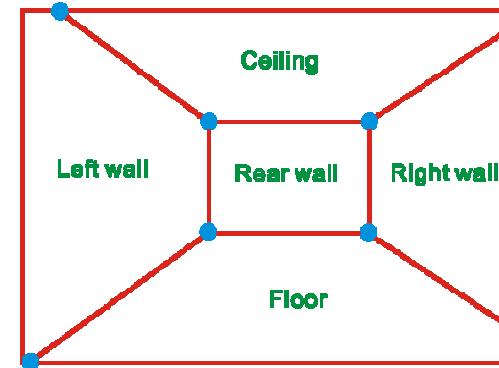


# Fun with vanishing points

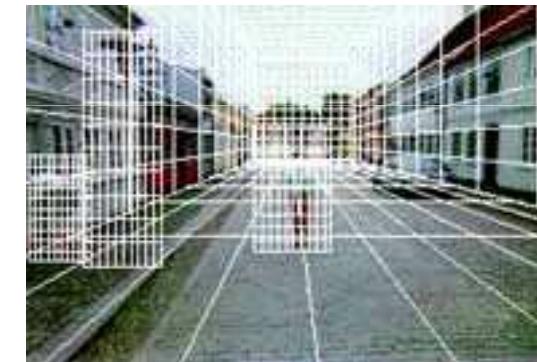


# “Tour into the Picture” (SIGGRAPH '97)

Create a 3D “theatre stage”  
of five billboards



Specify foreground objects  
through bounding polygons



Use camera transformations  
to navigate through the  
scene



# The idea

Many scenes (especially paintings), can be represented as an axis-aligned box volume (i.e. a stage)

Key assumptions:

- All walls of volume are orthogonal
- Camera view plane is parallel to back of volume
- Camera up is normal to volume bottom

How many vanishing points does the box have?

- Three, but two at infinity
- Single-point perspective

Can use the vanishing point  
to fit the box to the particular  
Scene!

# The idea

Many scenes (especially paintings), can be represented as an axis-aligned box volume (i.e. a stage)

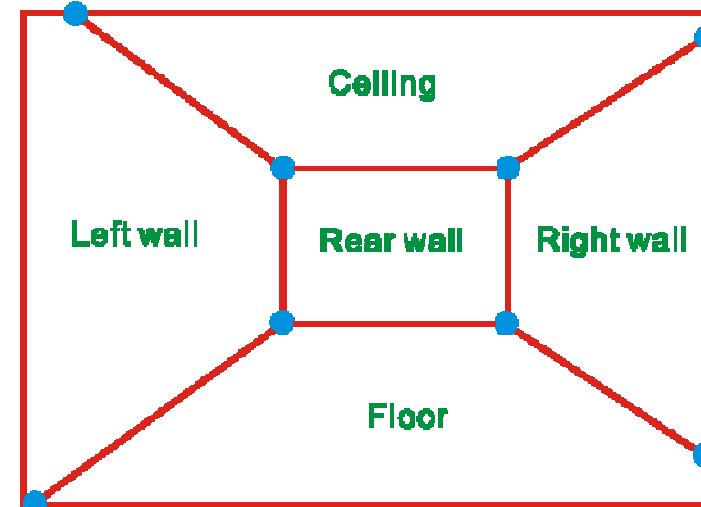
Key assumptions:

- All walls of volume are orthogonal
- Camera view plane is parallel to back of volume
- Camera up is normal to volume bottom

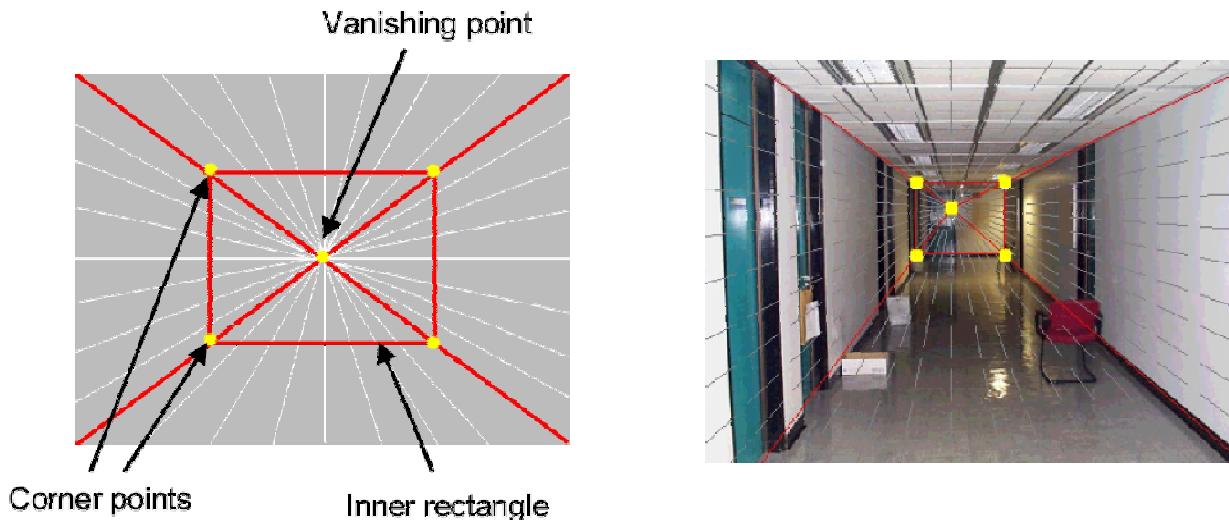
How many vanishing points does the box have?

- Three, but two at infinity
- Single-point perspective

Can use the vanishing point  
to fit the box to the particular  
Scene!



# Fitting the box volume

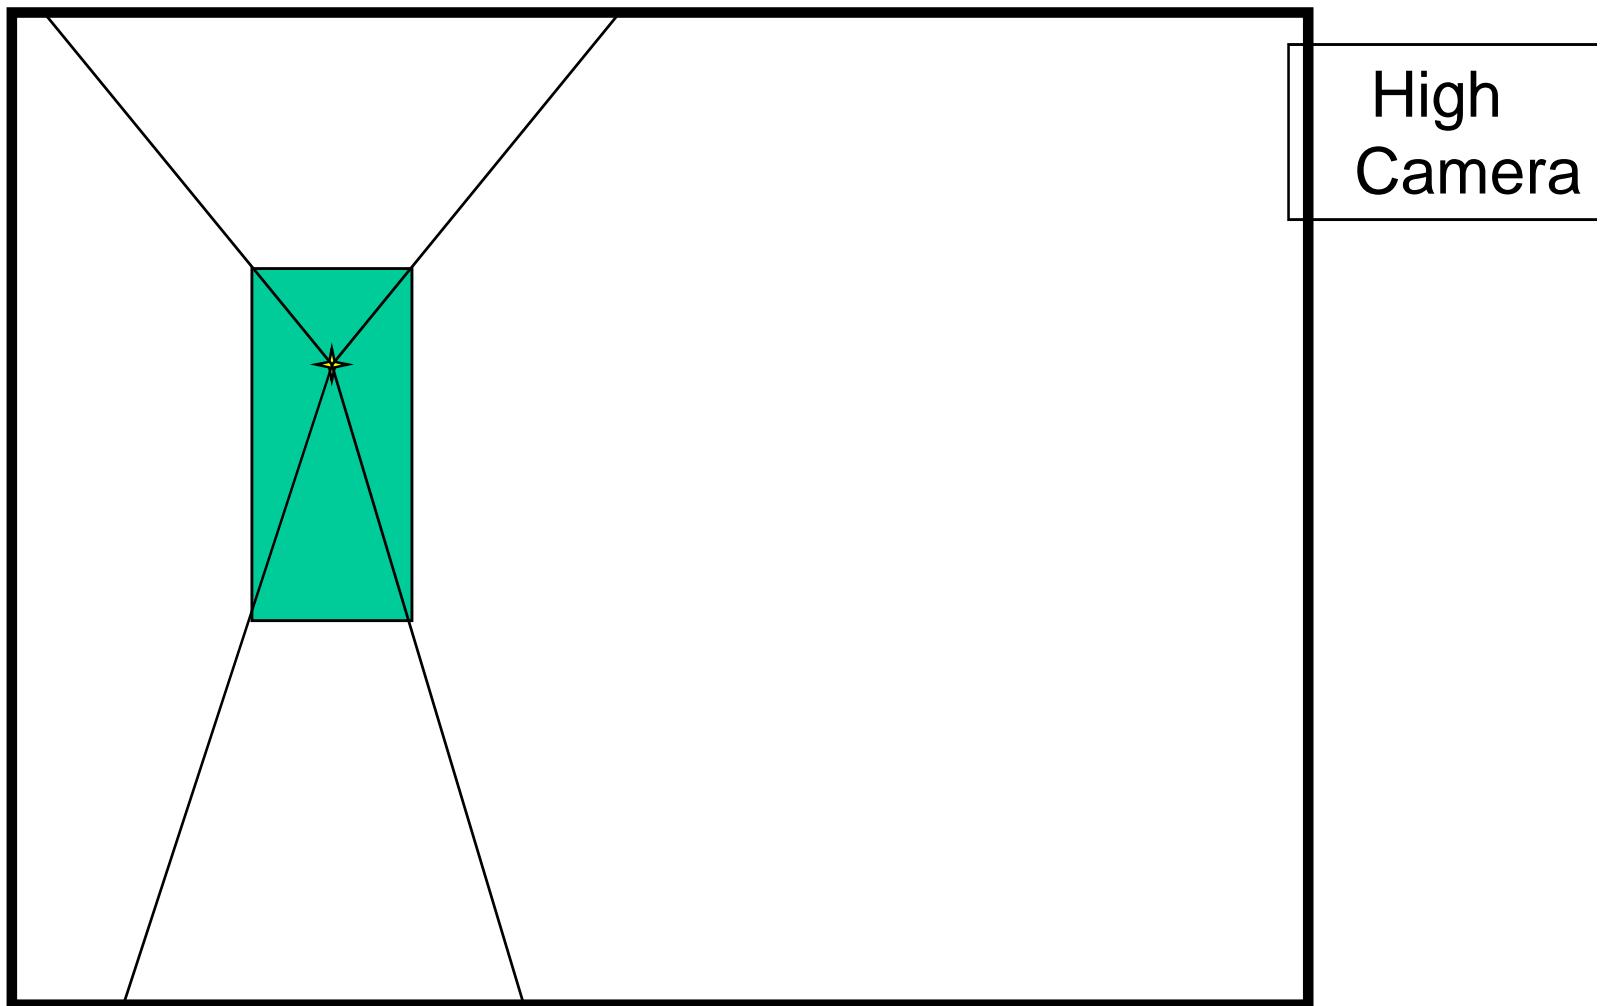


User controls the inner box and the vanishing point placement (# of DOF???)

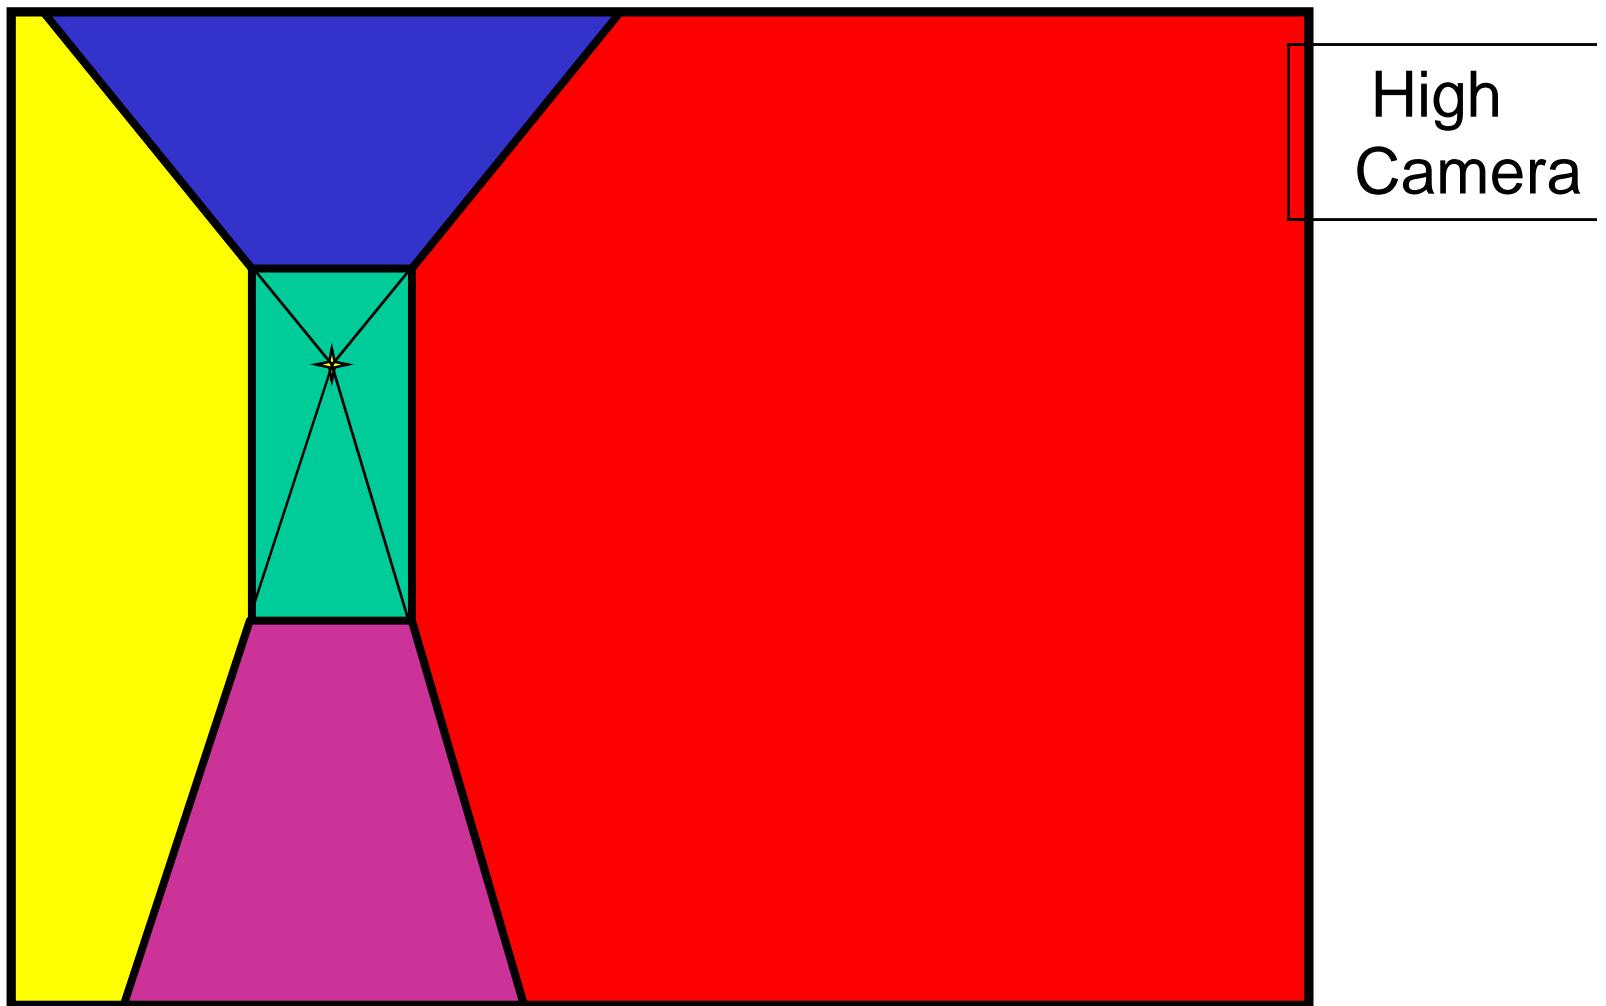
Q: What's the significance of the vanishing point location?

A: It's at eye level: ray from COP to VP is perpendicular to image plane.

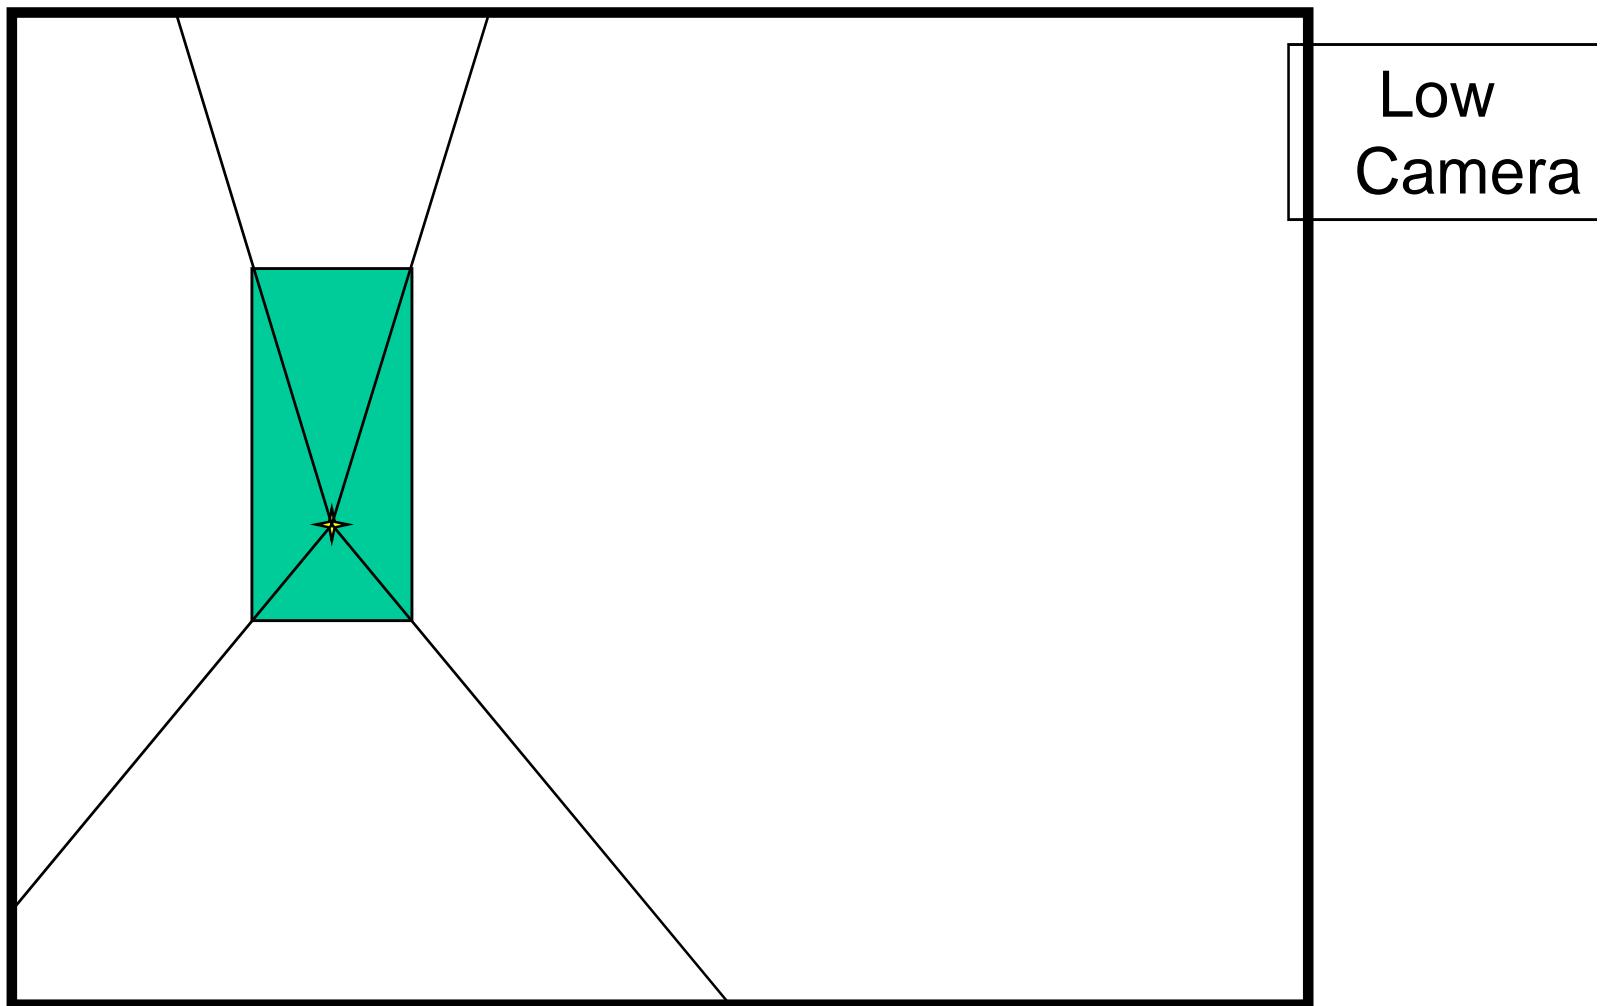
Example of user input: vanishing point and  
back face of view volume are defined



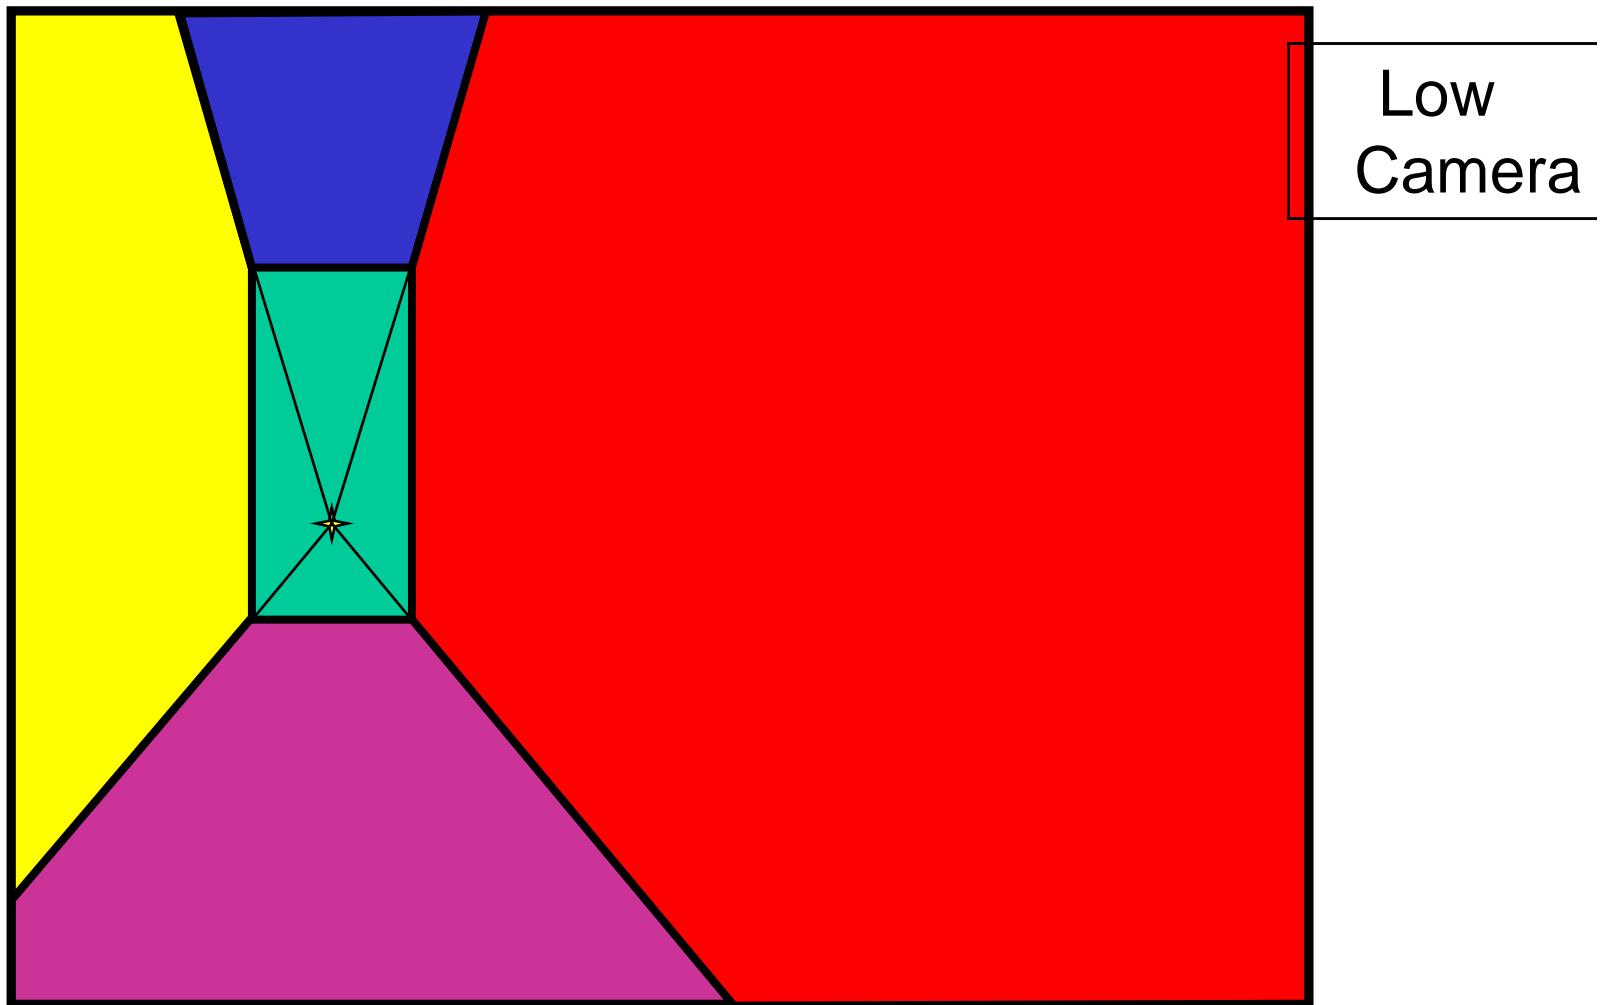
Example of user input: vanishing point and  
back face of view volume are defined



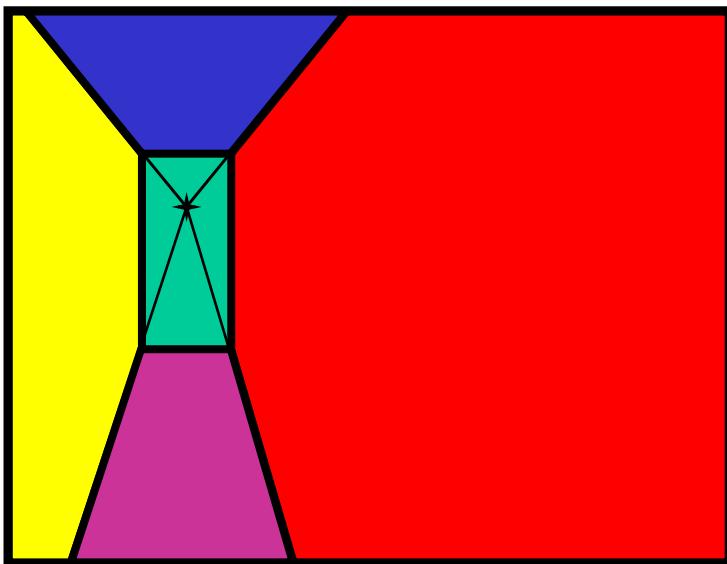
Example of user input: vanishing point and  
back face of view volume are defined



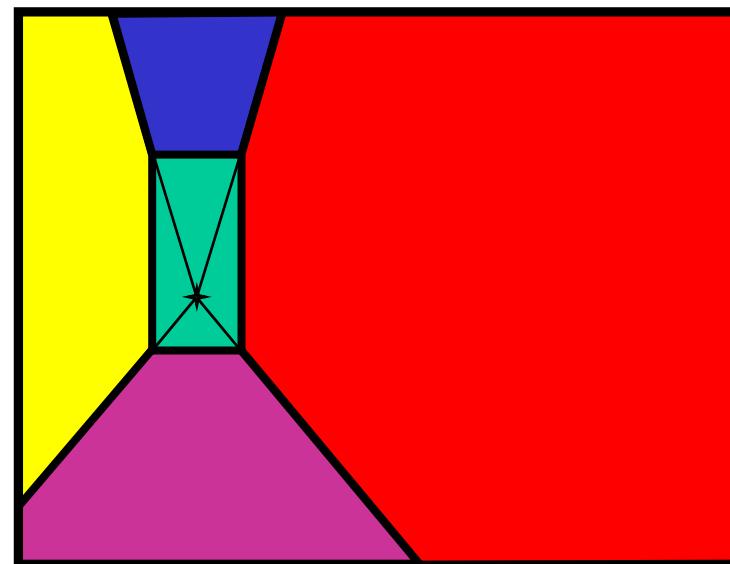
Example of user input: vanishing point and  
back face of view volume are defined



Comparison of how image is subdivided based on two different camera positions. You should see how moving the vanishing point corresponds to moving the eyepoint in the 3D world.

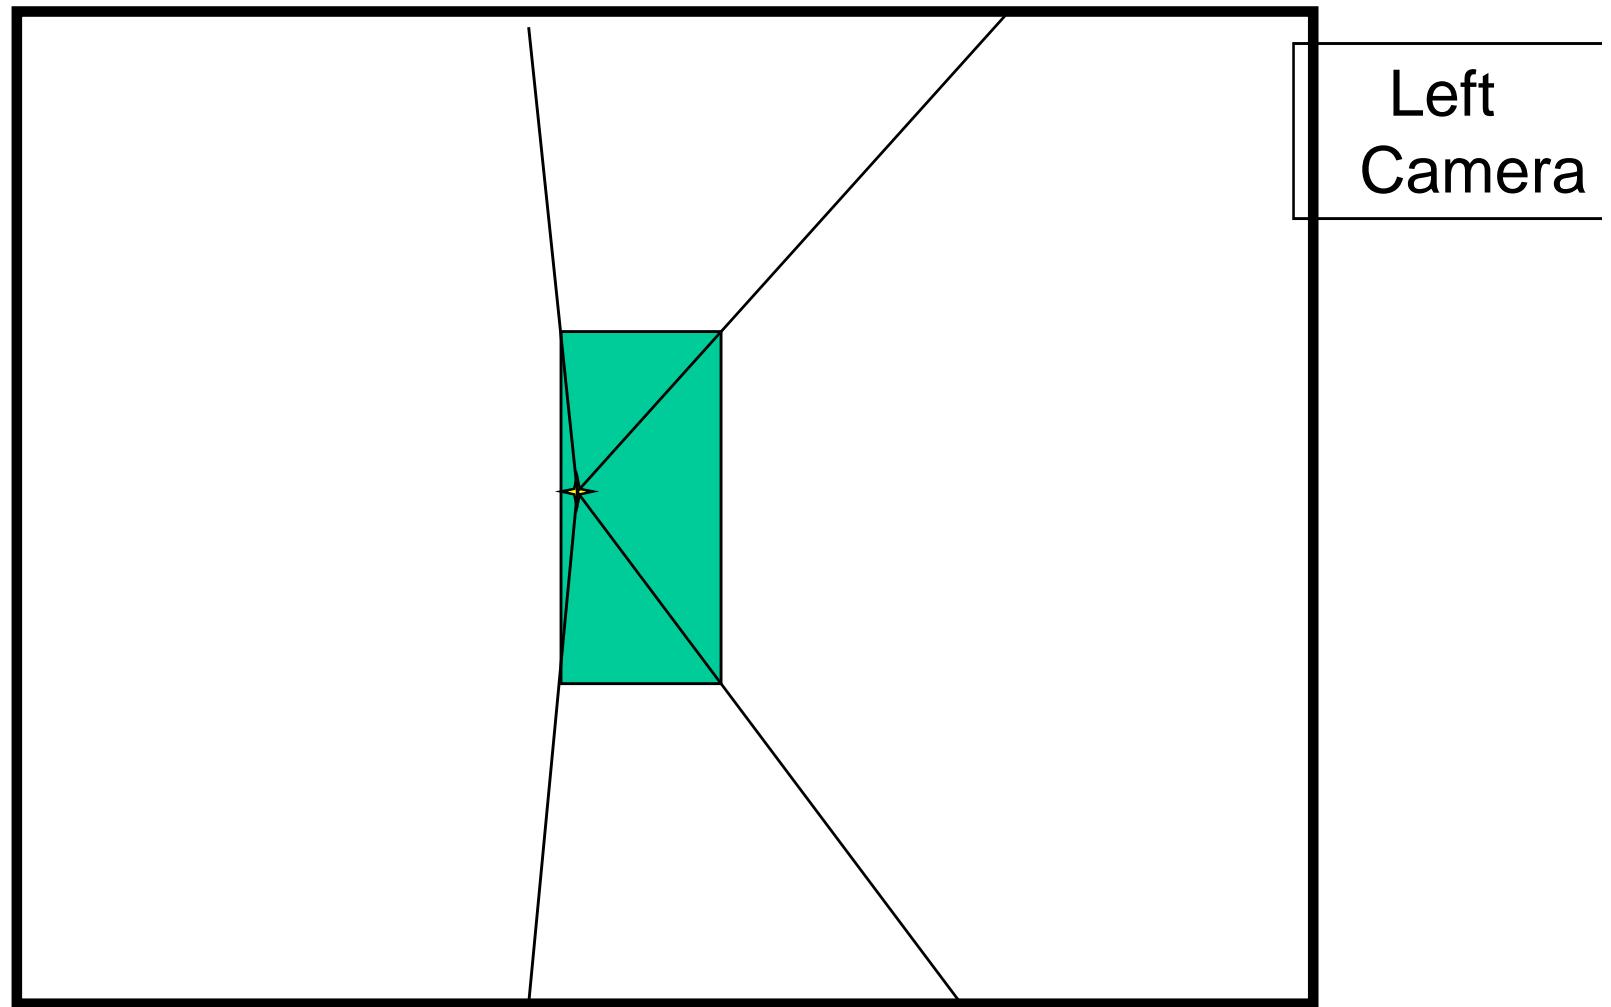


High Camera

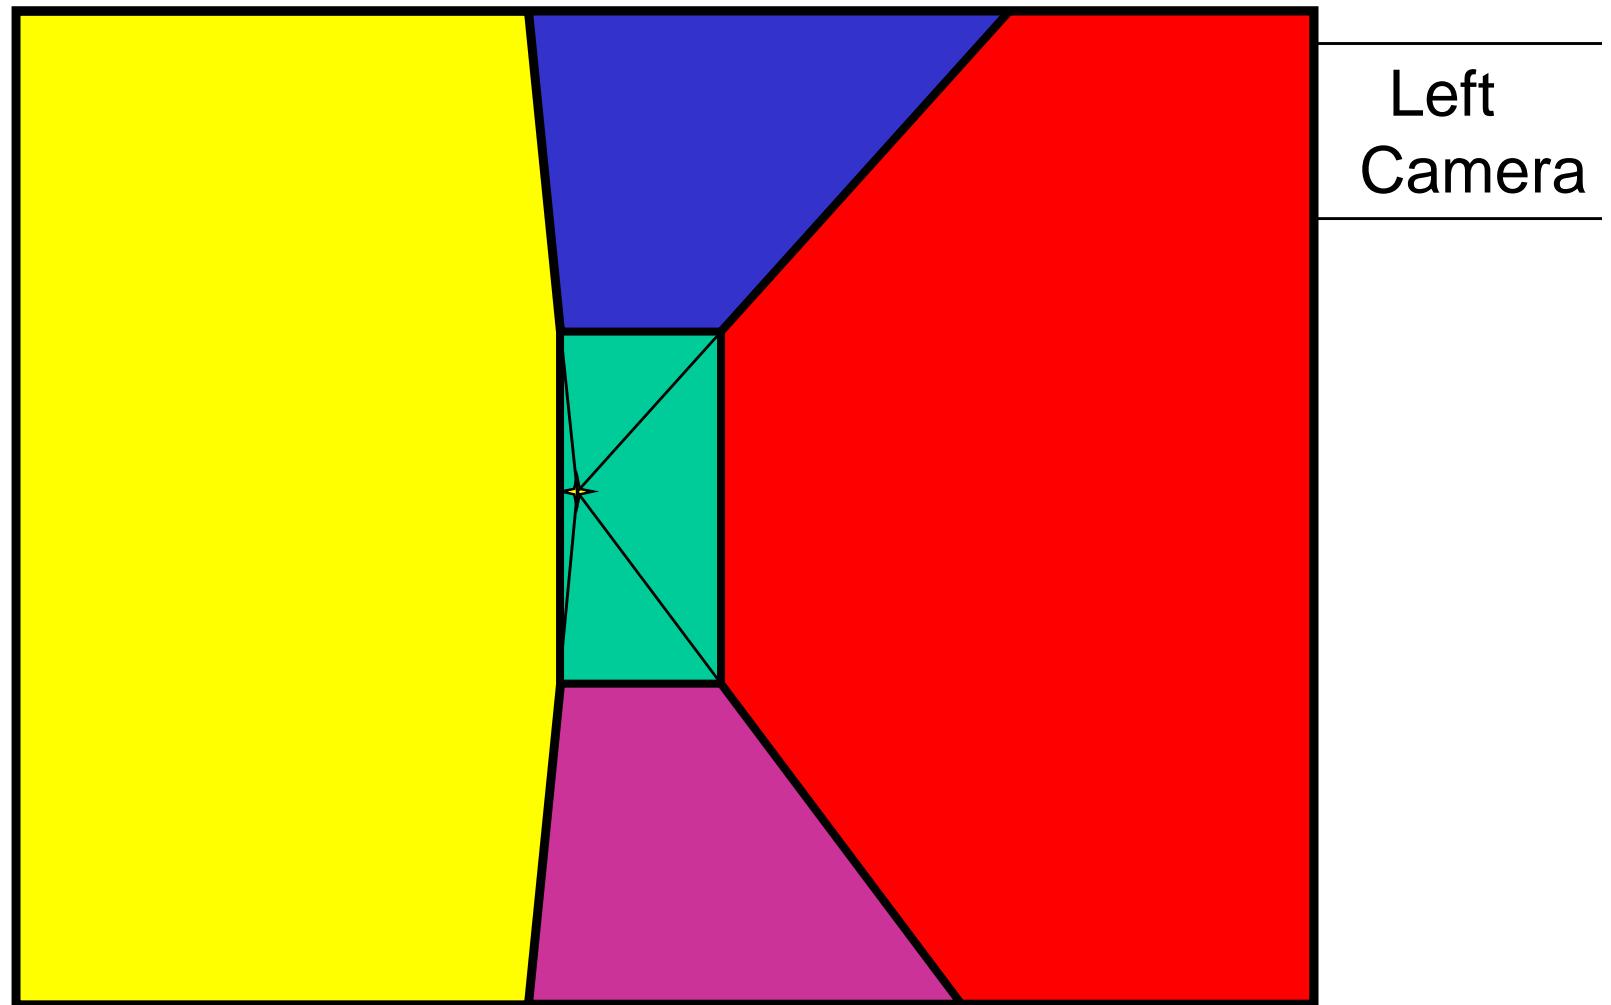


Low Camera

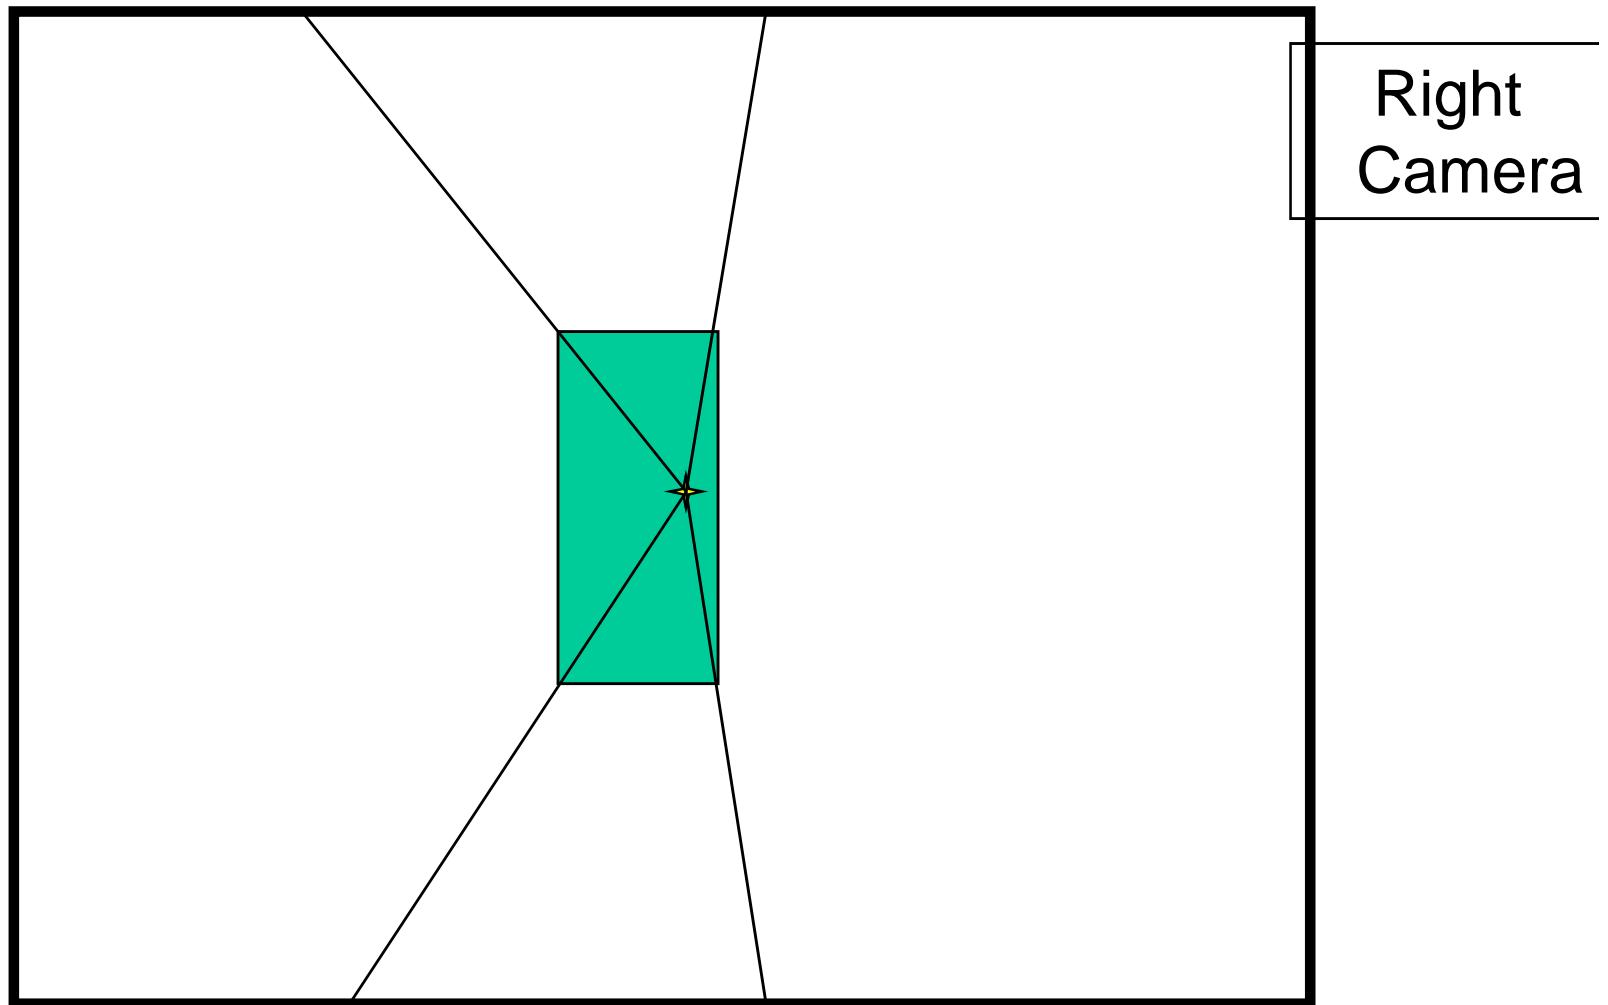
Another example of user input: vanishing point and back face of view volume are defined



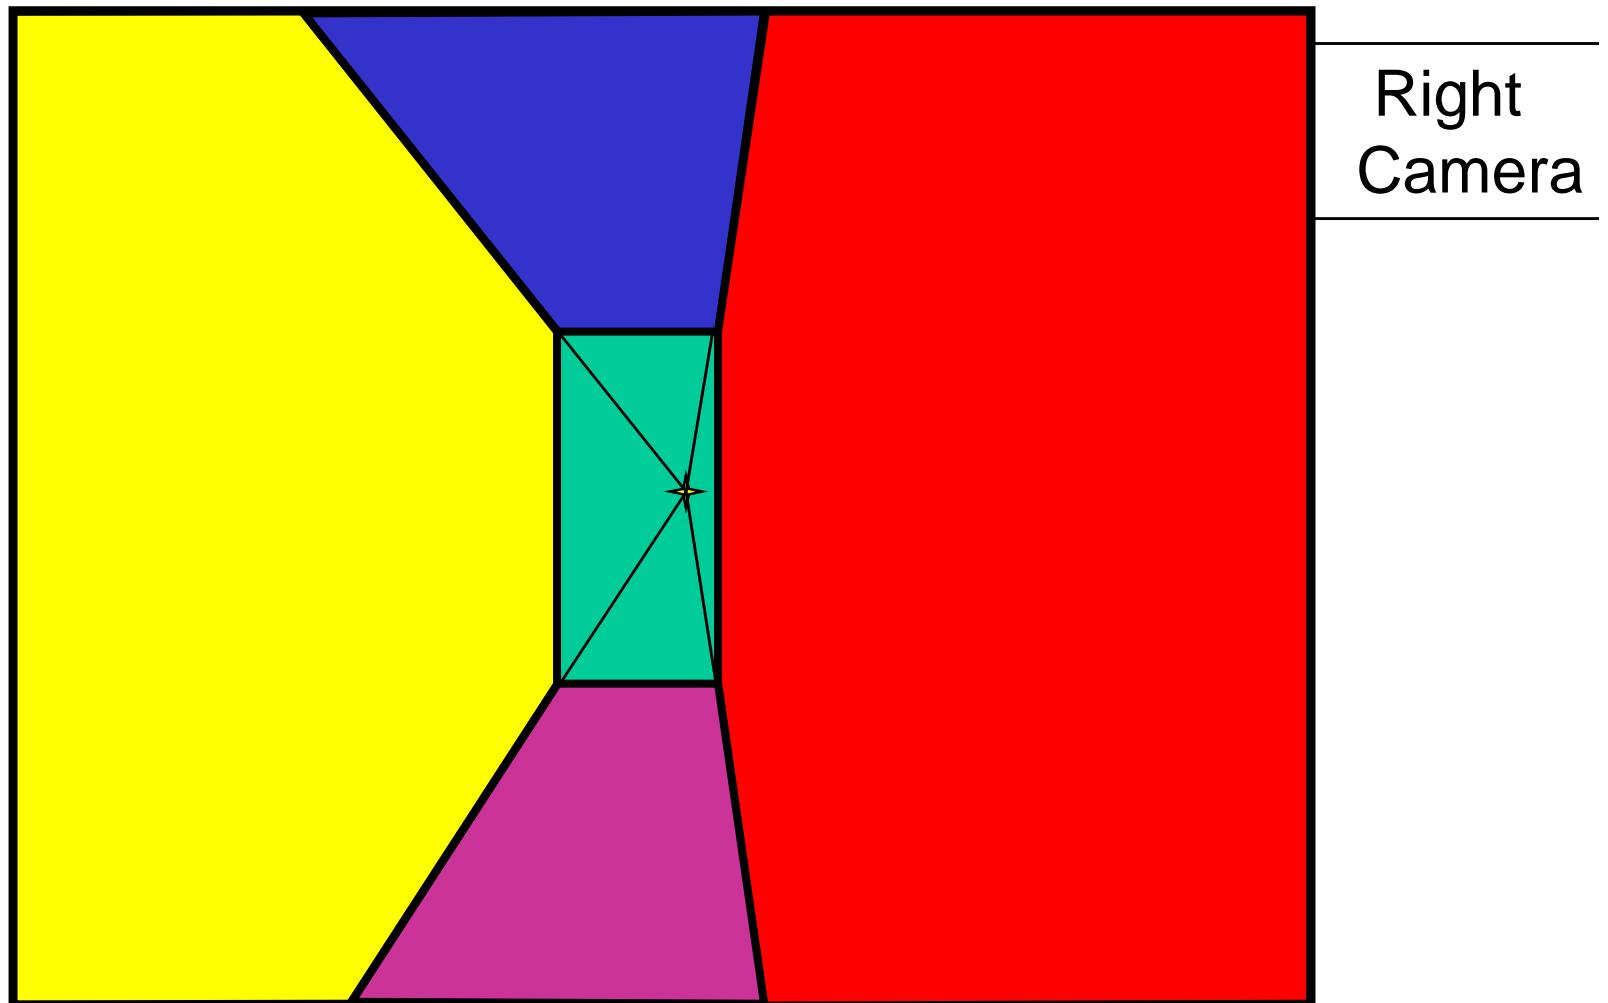
Another example of user input: vanishing point and back face of view volume are defined



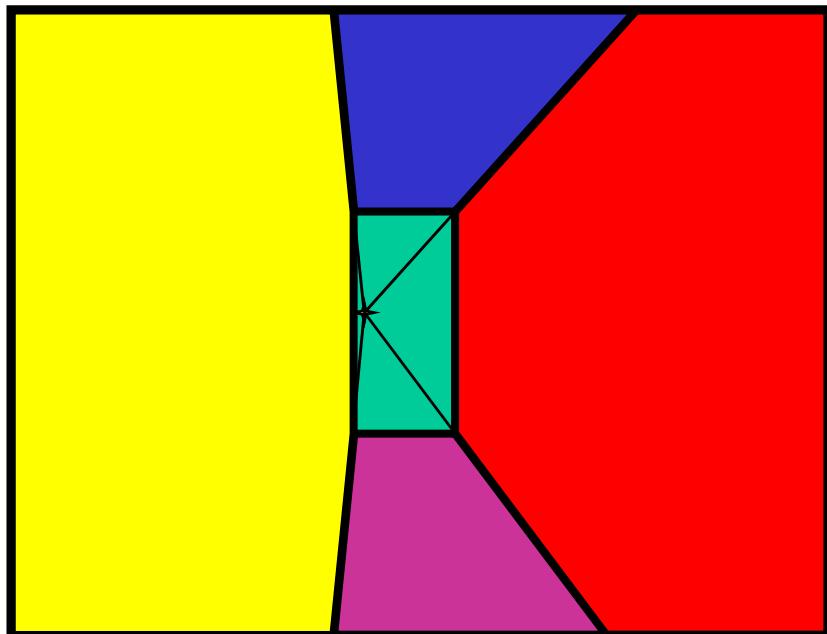
Another example of user input: vanishing point and back face of view volume are defined



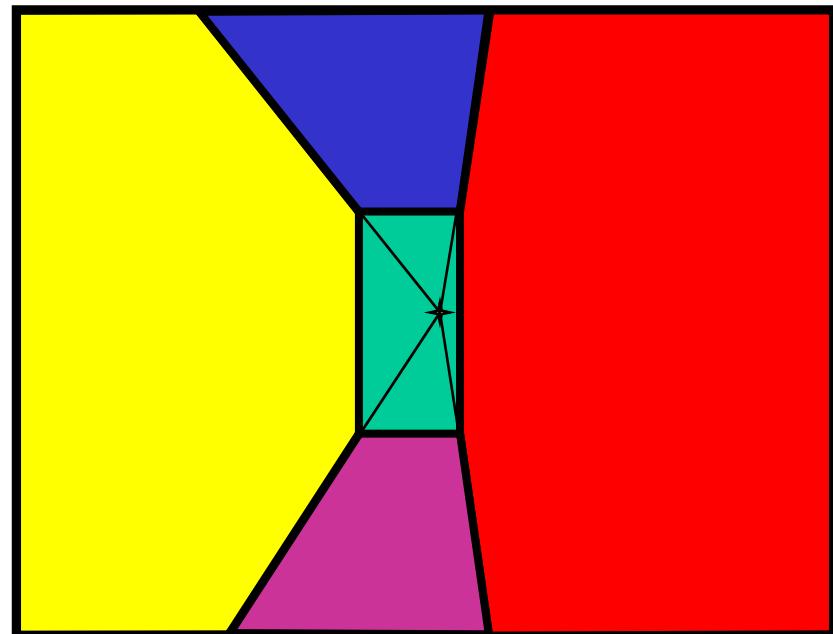
Another example of user input: vanishing point and back face of view volume are defined



Comparison of two camera placements – left and right.  
Corresponding subdivisions match view you would see if  
you looked down a hallway.



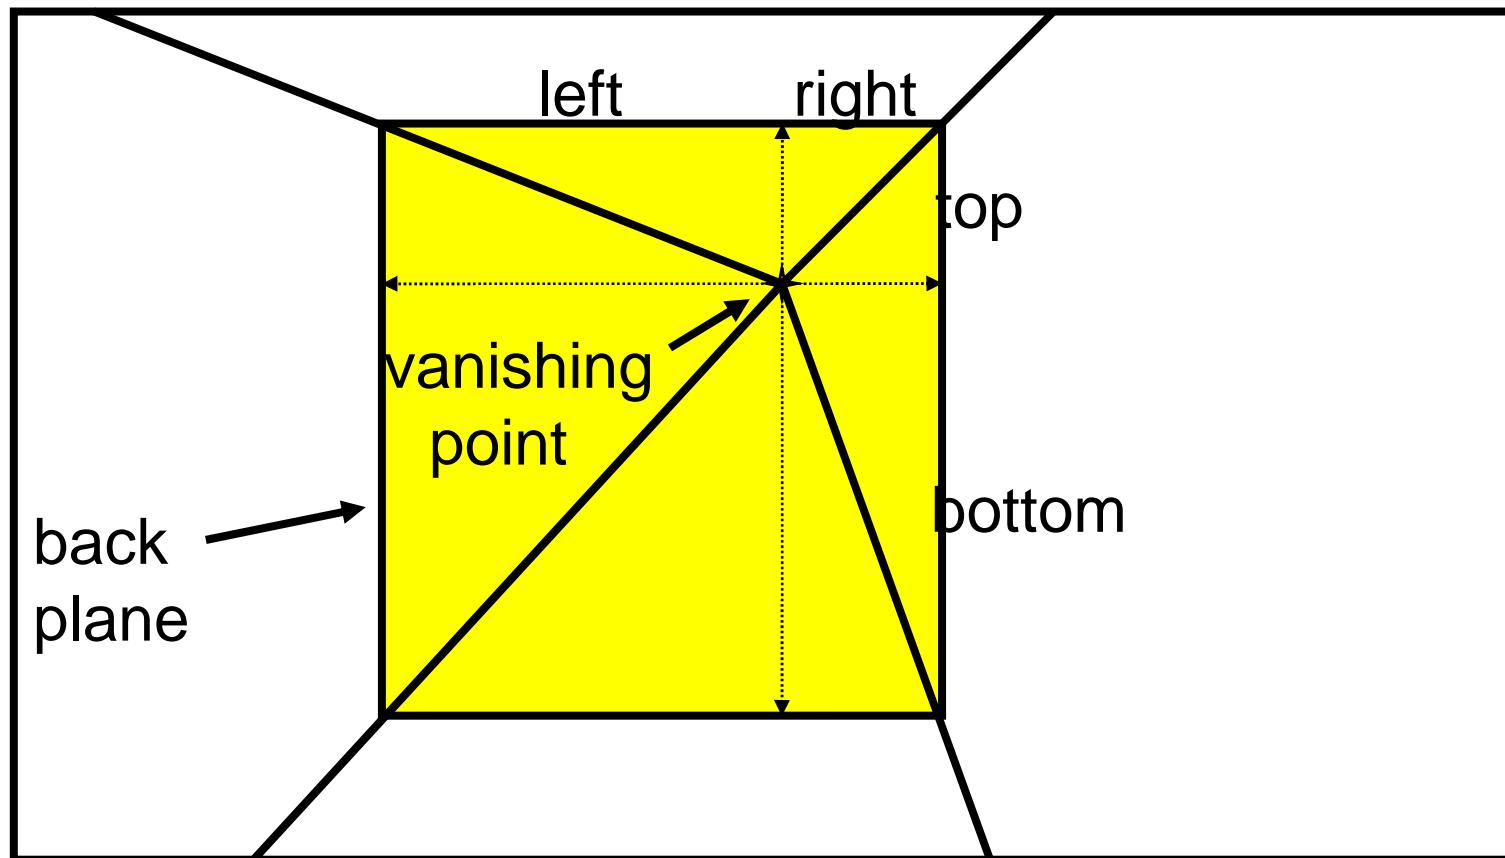
Left Camera



Right Camera

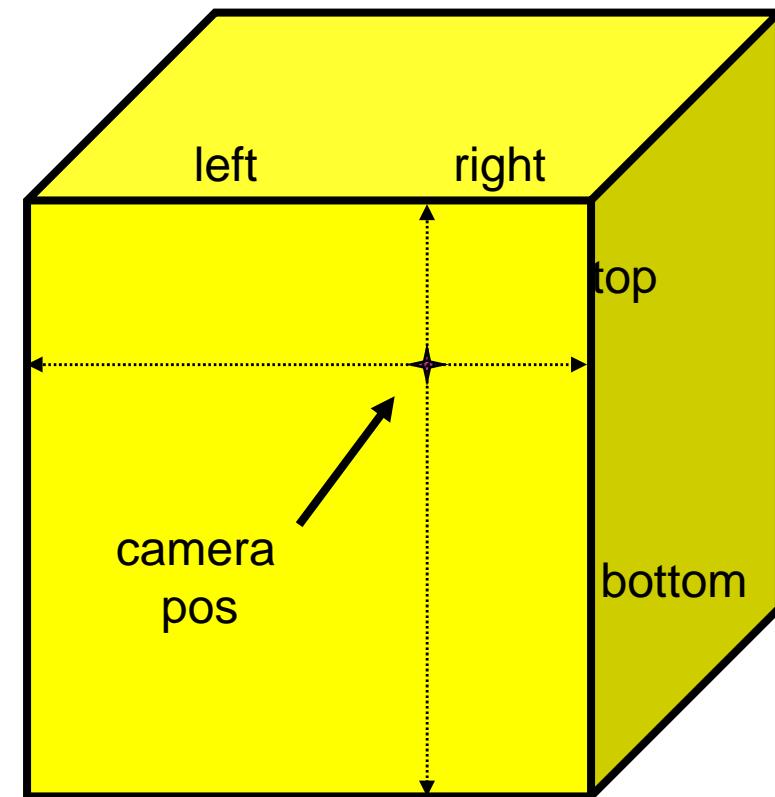
# 2D to 3D conversion

First, we can get ratios

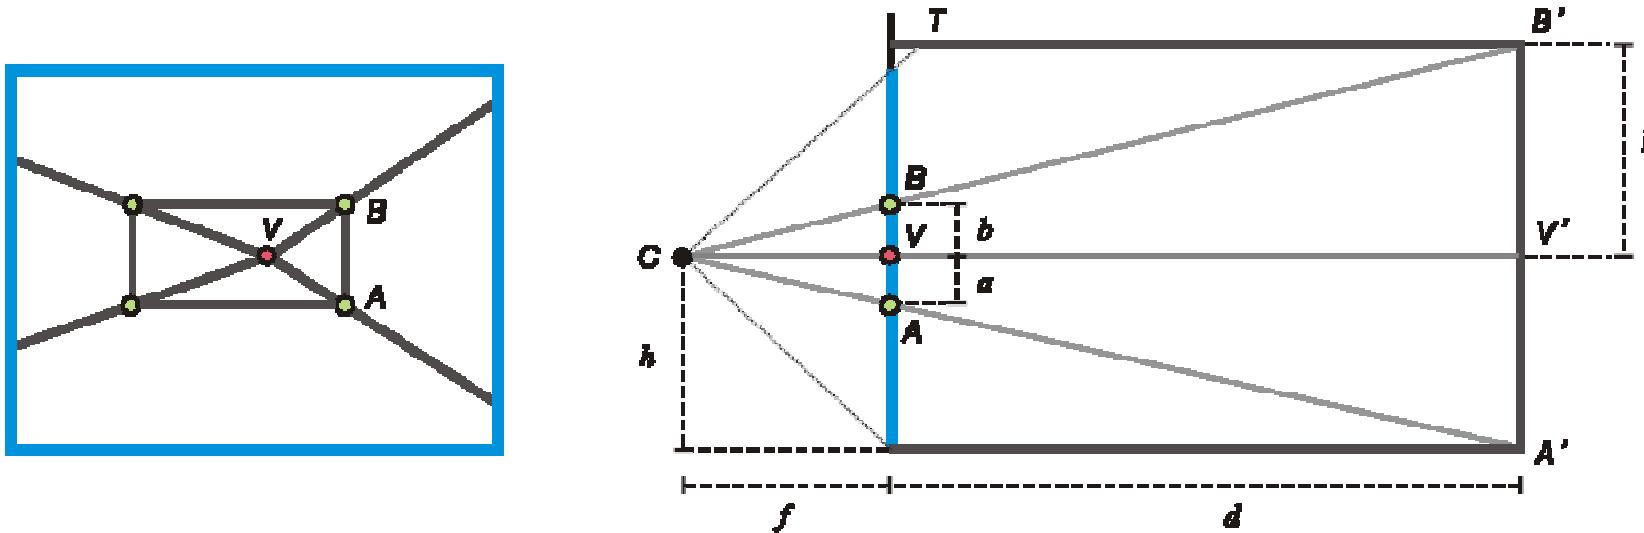


# 2D to 3D conversion

- Size of user-defined back plane must equal size of camera plane (orthogonal sides)
- Use top versus side ratio to determine relative height and width dimensions of box
- Left/right and top/bot ratios determine part of 3D camera placement



# Depth of the box



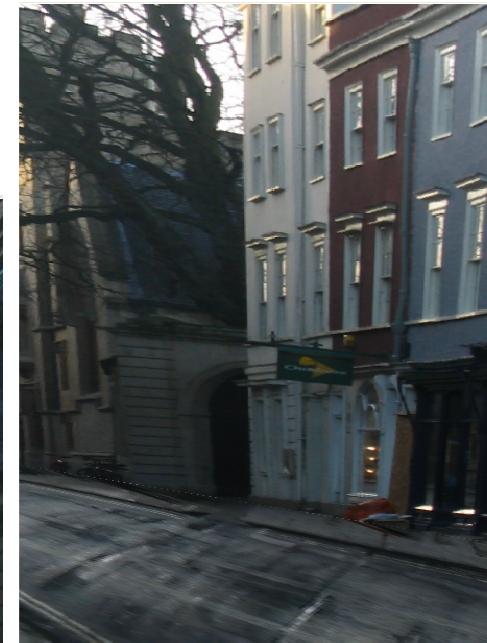
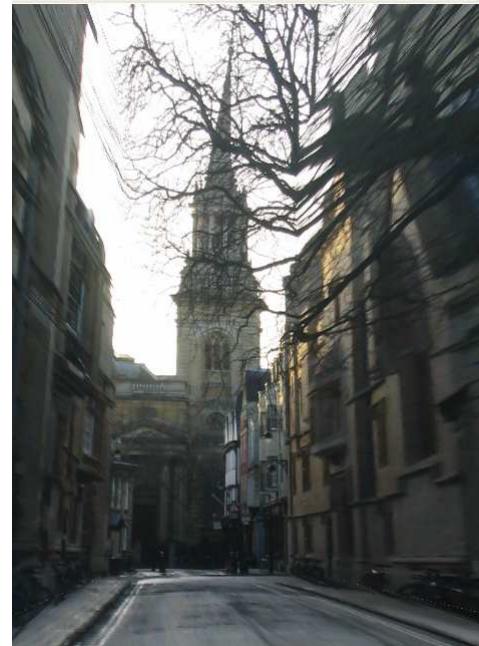
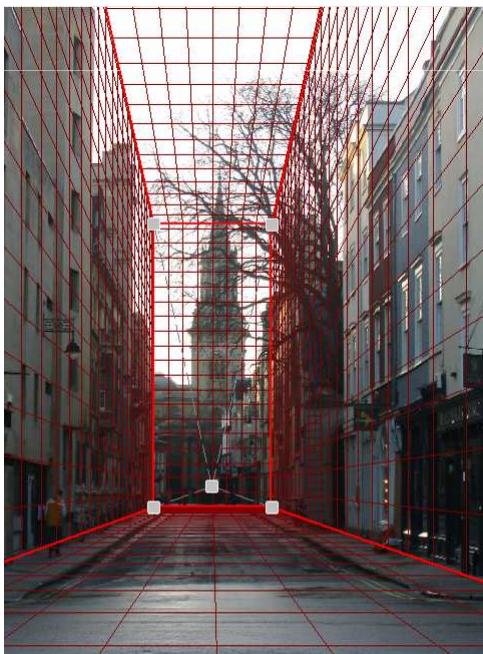
Can compute by similar triangles (CVA vs. CV'A')  
Need to know focal length  $f$  (or FOV)

Note: can compute position on any object on the ground

- Simple unprojection
- What about things off the ground?

# DEMO

Now, we know the 3D geometry of the box  
We can texture-map the box walls with texture from the image

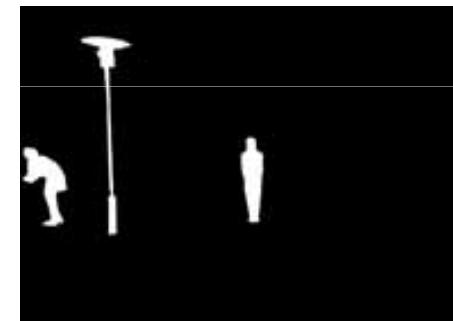


# Foreground Objects

Use separate billboard for each

For this to work, three separate images used:

- Original image.
- Mask to isolate desired foreground images.
- Background with objects removed

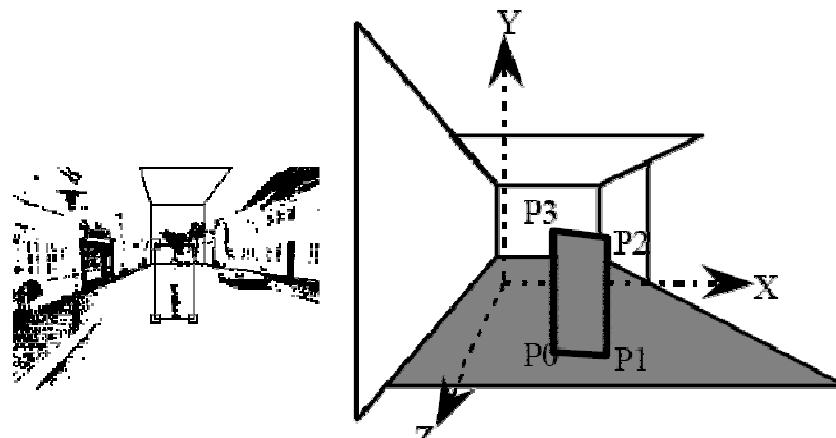


# Foreground Objects

Add vertical rectangles  
for each foreground  
object

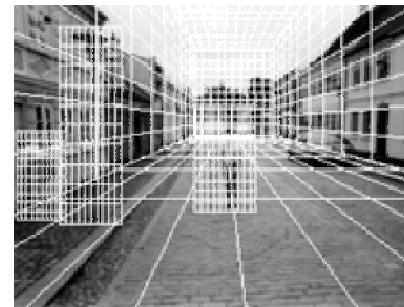
Can compute 3D  
coordinates P0, P1 since  
they are on known plane.

P2, P3 can be computed  
as before (similar  
triangles)



(a) Specifying of a foreground object

(b) Estimating the vertices of the foreground object model



(c) Three foreground object models

# Foreground DEMO

